

Governance Concentration Beyond Token Allocation: A 40-Protocol Cross-Sectional Audit of DePIN and DeFi After Protocol-Controlled-Address Exclusion

Zach Zukowski

Tokenization Systems

zach@tokenization.systems

ORCID: 0009-0006-3642-2450

April 2026

Abstract

Initial token allocation design is not associated with post-distribution governance concentration. This paper reports a 40-protocol cross-section spanning DePIN, DeFi, infrastructure, and social token categories, with protocol-controlled address (PCA) exclusions applied before computing Herfindahl-Hirschman Index (HHI) concentration metrics. Team and investor allocation at launch is uninformative about steady-state concentration (Pearson $r = 0.05$, $p = 0.76$, $N = 37$). Post-distribution dynamics dominate launch design: protocols with more insider wallets among top holders show higher concentration (Spearman $\rho = 0.48$, $p = 0.003$, $N = 37$), and the association survives a non-insider HHI tautology check ($\rho = 0.54$, $p = 0.001$, $N = 34$).

Three further findings sharpen the measurement picture. DePIN governance is more concentrated than DeFi after PCA correction (Mann-Whitney $p = 0.020$, Cohen's $d = 0.94$; robust in 30/30 leave-one-out iterations). Subsidy intensity is associated with concentration only when Livepeer is included ($r = 0.59$, $p = 0.003$, $N = 23$); excluding that single outlier yields a clean null ($r = 0.07$, $p = 0.77$). Delegation amplifies voting power above token holdings in nine of thirteen protocols with sufficient governance data (ratios 1.6x to 9.9x; mean approximately 4.9x), with four design-driven exceptions (ENS, GMX, HNT, JUP).

The methodological contribution is a generalizable PCA exclusion typology correcting systematic inflation in prior holder-list concentration studies (median inflation factor 2.3x; maximum approximately 18x). Gini and HHI capture distinct properties of the same holder set ($r = 0.59$); inequality metrics cannot substitute for direct concentration measurement.

JEL Codes: C21, C81, D02, D71, D72, G23, G34, L22, O33

Keywords: governance concentration, Herfindahl-Hirschman Index, DePIN, DAO governance, delegation amplification, protocol-controlled addresses, token allocation, measurement methodology

1 Introduction

1.1 Motivation

Token-based governance is the default coordination mechanism for billions of dollars in decentralized physical infrastructure (DePIN) and decentralized finance (DeFi). But what shapes whether governance power concentrates or distributes? Three candidate explanations dominate practitioner discourse: insider token allocation concentrates power; subsidy dependence (high token emissions relative to revenue) entrenches early holders; and younger protocols have not yet had time to decentralize. This paper tests all three and finds that none is sufficient as a standalone explanation of the cross-sectional variation observed across forty protocols.

Two empirical contrasts illustrate why the candidate explanations fail. Both ENS and DIMO let token holders delegate their voting power, yet the two protocols produce opposite governance concentration outcomes. ENS (the Ethereum Name Service) funds roughly 100 active community delegates, and the resulting voting-power distribution is less concentrated than the underlying token holdings (voting-to-holding amplification ratio 0.45x; ratios below 1.0 mean delegation disperses governance power relative to the holding baseline). DIMO, a vehicle-data DePIN protocol, has ten active voters across the most recent twelve months, and its voting-power distribution is 9.1 times more concentrated than its token holdings (amplification ratio 9.1x). Across the thirteen-protocol governance sample studied here, delegation amplifies concentration in nine of thirteen protocols (range 1.6x to 9.9x above the holding baseline), with four structural exceptions on the dispersing side (ENS 0.45x, GMX 0.87x, HNT 0.35x to 0.53x, JUP 0.057x; Section 4.5). The cross-protocol variation is not associated with sector membership or governance venue but by institutional design choices at the delegate-program level: how many community delegates the protocol funds, how voting power flows from holders to delegates, and what incentives shape delegate participation (Section 5.1).

A parallel pattern emerges for subsidy economics: financial sustainability is uninformative as a concentration indicator. Three protocols with nearly identical subsidy ratios (Aethir 0.36, Aave 0.37, io.net 0.40) span nearly an order of magnitude in governance HHI (0.013 to 0.125): financial sustainability is neither necessary nor sufficient for governance decentralization (Section 4.4). Governance concentration emerges as a secondary market outcome, shaped by post-distribution trading and accumulation dynamics that neither allocation design nor protocol financial characteristics explain in this cross-section.

We construct a 40-protocol governance concentration dataset spanning DePIN and DeFi, develop a Blockscout-verified (a blockchain contract verification tool) exclusion methodology to remove 133 protocol-controlled address exclusions (PCAs; addresses that appear on holder lists but structurally cannot vote, including staking contracts that hold tokens on behalf of stakers, exchange custodians that custody customer deposits, vesting locks that release tokens to insiders over time, and protocol treasuries) from HHI computation, and test every available cross-sectional covariate of concentration using both parametric and rank-based methods.

Four empirical contributions emerge. First, naive token holder lists include addresses that cannot vote: staking contracts, exchange custodians, vesting locks, and protocol treasuries. Across 38 protocols, 133 such address inclusions (125 unique addresses; five recur across multiple protocols, e.g., Binance 8 hot wallet across four tokens) inflated affected protocols' HHI in prior

concentration measurements by up to 18x (RENDER, where Wormhole Token Bridge custody dominated naive top-1000 holdings; median inflation factor 2.3x across 32 protocols with complete pre-exclusion and post-exclusion HHI logged). The methodology generalizes via the five-class typology in Section 3.8, supporting uniform replication across ERC-20 and SPL token-governance datasets. Second, initial allocation design is uninformative about concentration while current insider wallet retention is associated with concentration (Spearman $\rho = 0.48$, $p = 0.003$, $N = 37$; Section 4.4). Third, DePIN governance concentration is higher than DeFi after exclusion correction (Mann-Whitney $p = 0.020$, Cohen's $d = 0.94$), and the finding is robust to single-protocol exclusion (significant in all 30 of 30 leave-one-out iterations).

Fourth, delegation concentrates governance power in nine of thirteen protocols sampled, with amplification factors of 1.6x (DRIFT), 2.7x (Uniswap), 3.1x (Arbitrum), 3.6x (Optimism), 3.8x (WeatherXM), 4.4x (Aave), 6.3x (Lido), 9.1x (DIMO), and 9.9x (Compound). Four protocols (ENS at 0.45x; GMX at 0.87x; HNT at 0.35x-0.53x; JUP at 0.057x most-extreme dispersion outlier) show delegation-mediated dispersion, reflecting mature-delegate-program (ENS), methodology-sensitive crossover (GMX), Solana VSR lockup-weighted design (HNT), and Solana delegated-vote with broad airdrop user base (JUP) patterns where holders systematically delegate to a broader community-delegate set or vote-weight mechanism design distributes voting power below the holding baseline. Two further protocols (Uniswap, Optimism) would appear to show dispersion if holding HHI were measured without excluding addresses that cannot vote (burn wallets, foundation treasuries); under the consistent exclusion methodology applied throughout this paper, both protocols join the amplification pattern. Curve and Balancer (vote-escrowed governance via veCRV, which is vote-escrowed CRV, and veBAL, vote-escrowed BAL) show extreme amplification (15x and 21x) reflecting lock-duration weighting and form a distinct class discussed in Section 4.5.2. Institutional design within each delegate program is associated with whether delegation concentrates power and how strongly.

Recent scoping work by Alshater (2026) systematically maps the DePIN tokenomics design space, identifying the burn-and-mint equilibrium, governance-calibrated issuance, and provider staking as the dominant economic primitives across leading deployments. Alshater explicitly frames decentralization as a spectrum rather than a binary property, echoing analyses of decentralized governance for cyber-physical systems (Nabben et al., 2024), and identifies the lack of rigorous quantitative empirical work on these models as the field's most pressing research gap. The cross-sectional measurement reported here responds to that gap, applying Herfindahl-Hirschman Index analysis to forty protocols spanning DePIN and adjacent categories.

These results extend the governance concentration literature (Barbereau et al., 2023; Feichtinger et al., 2023; Cong et al., 2025) through a methodology that corrects systematic measurement artifacts in prior holder-list studies and through cross-sectional evidence on allocation, sector, subsidy, and delegation covariates of concentration.

1.2 Companion Theory and Normative Bridge

Governance concentration is not only a statistical pattern; it bears on institutional legitimacy when a small set of holders or delegates can set rules affecting many participants. A companion theoretical paper develops the Polanyian interpretation: governance rights packaged as transferable tokens function as a candidate fourth fictitious commodity, and concentration patterns are consistent with commodification of constitutive political authority rather than

transient implementation failure (Zukowski, 2026e). This empirical paper does not develop that framework at length. It supplies the cross-sectional measurement base Zukowski (2026e) cites for Test 2 and documents where delegate-program design, not sector or launch allocation, shapes voting-power outcomes.

1.3 Research Questions

Three empirical research questions guide the analysis:

RQ1: After PCA exclusion, which cross-sectional protocol features are associated with post-distribution holder concentration (HHI)?

RQ2: How does delegated voting power deviate from post-exclusion token holdings, and is amplification associated with sector or governance venue?

RQ3: Which measurement and design implications follow for governance audits that rely on launch-time allocation or naive holder lists?

1.4 Scope, Assumptions, and Limitations

1.4.1 Scope

The empirical analysis focuses on tokenized systems between 2017 and 2025 across DeFi, infrastructure protocols, DAOs, and DePIN networks. The core analytic target is **governance concentration measurement** and its associations with protocol characteristics, not short-term market performance.

1.4.2 Assumptions

- **Holder lists require classification.** Addresses that custody, burn, lock, bridge, or treasury-hold tokens should not be treated as independent governance holders.
- **Snapshots measure states, not trajectories.** The March 2026 cross-section supports association and comparison, not causal claims about deconcentration over time.
- **Voting power can diverge from holdings.** Delegation and vote-escrow mechanics require separate voting-HHI measurement where data are available.

1.4.3 Limitations (preview)

Three classes of limitation deserve early notice. Measurement scope: concentration is measured from holder lists and governance voting surfaces after PCA exclusion; informal influence, off-chain coordination, and proposal-specific quorum rules are not fully captured. Causal inference scope: a March 2026 cross-section supports descriptive association, not causal identification. Data limitations: voting-HHI is available for thirteen protocols, PCA classification is weaker on some Solana-native tokens, and the 40-protocol sample is small for fully specified multivariate models. Each limitation is developed in Section 5.7 with the relevant robustness defenses.

1.5 Paper Outline

Section 2 reviews related work. Section 3 specifies data, PCA exclusion methodology, and statistical tests. Section 4 presents results. Section 5 discusses design implications and

limitations. Section 6 concludes. Supplementary materials provide replication scripts, robustness analyses, and the canonical statistics ledger (Supplementary File S0).

2 Related Work

Sections 2.1 through 2.8 synthesize related literature on token governance, institutional design, DePIN, and concentration measurement. Section 3 specifies the empirical pipeline, PCA exclusion methodology, and measurement instruments used in Sections 4 and 5. Readers focused on the empirical contribution may read Section 3 first.

2.1 Overview

Three measurement gaps in the token-governance literature motivate this paper. First, prior holder-list HHI computations have not systematically excluded protocol-controlled addresses (staking contracts, treasuries, bridge custody, exchange wallets), inflating concentration estimates by a median 2.3x and up to 18x in this paper’s sample. Second, the cross-sectional relationship between launch-time token allocation and post-distribution concentration has not been tested with sample size sufficient for null findings to be informative. Third, the gap between holding HHI and delegated voting HHI has not been measured across sectors using comparable governance data surfaces (Tally, Snapshot, Solana on-chain governance).

The literature spans voting-power concentration in major DAOs (Fritsch et al., 2024), insider trading by proposal managers (Cong et al., 2025), the “decentralization illusion” (Aramonte et al., 2021), and structural concentration patterns across categories (Barbereau et al., 2023; Feichtinger et al., 2023). The subsections below review the strands relevant to each gap: tokenomics as mechanism and institutional design (§2.2-2.3), DAO governance concentration and delegation (§2.4), DePIN tokenomics (§2.5), constitutional and polycentric framings (§2.6), and synthesis (§2.7-§2.8).

2.2 Tokenomics and Cryptoeconomic Design

2.2.1 Tokenomics as a Mechanism Design Problem

Tokenomics can be understood as a form of mechanism design (Hurwicz, 1973; Myerson, 1981) and, more broadly, as an exercise in institutional invention (Hurwicz, 1987): the “reverse engineering” of game theory in which a designer specifies rules and incentives to induce desired behavior among strategic agents (users, validators, operators, token holders). Token-governance protocols are explicitly invented institutions, designed from whitepaper plus code rather than evolved from social practice, which sharpens the design question Hurwicz 1987 poses for invented institutions generally: any implementable mechanism must satisfy incentive compatibility (participants find truthful participation in their interest) and informational feasibility (the mechanism cannot require information no participant possesses). The mechanism design literature provides foundational formal tools for reasoning about incentive compatibility, equilibrium selection, and adverse selection. These tools remain highly relevant when tokenized systems must function under manipulation pressure and incomplete information. The Nobel Prize’s summary of mechanism design theory emphasizes incentive-compatible direct

mechanisms and the revelation principle as foundational concepts, offering a bridge between classical economics and modern crypto-economic system design.

Token systems, however, often introduce additional constraints absent from classical mechanism design (Roughgarden, 2021) settings: open entry, pseudonymity, composability, and adversarial execution environments. These constraints place a premium on robust incentive design, fraud resistance, and governance structures that can adapt without becoming arbitrary or captured.

2.2.2 Transaction Fee Mechanisms and “Sinks” as Economic Policy

Some of the clearest illustrations of tokenomics-as-policy appear in fee mechanism design. Ethereum’s EIP-1559 (Buterin et al., 2019; an Ethereum Improvement Proposal is a community-authored specification document defining a protocol change for Ethereum, analogous to a constitutional amendment for the system’s economic rules), for instance, introduced a base fee adjusted by network demand and burned rather than paid to miners/validators, creating a direct “sink” (a mechanism that permanently removes tokens from circulation, reducing supply and tying the token’s value to system usage) linked to usage and shaping the economic properties of the asset. Tim Roughgarden’s work on transaction fee mechanism design treats EIP-1559 not as marketing but as mechanism engineering: explicitly analyzing goals and trade-offs, including alternative designs such as smoothing and revenue-forwarding.

Subsequent research (Roughgarden, 2020) highlights that even well-designed sinks can introduce new strategic behaviors (e.g., manipulation or minority attacks on base fee dynamics), underscoring that tokenomics is not “set-and-forget” but a governance problem requiring transparency and monitoring.

These studies motivate a broader design intuition central to this paper: token mechanics that tie usage to scarcity (through burns, locks, buybacks, or fee routing) represent a form of fiscal policy embedded in code (cf. Werner et al., 2022).

The leading formal model of token adoption (Cong, Li, & Wang, 2021) shows that equilibrium token price is determined by aggregating heterogeneous users’ transactional demand rather than by initial supply distribution, implying that post-distribution dynamics dominate design-time allocation choices. This theoretical result anticipates the allocation design null documented in Section 4.4.

2.3 Incentives, Liquidity Mining, and the “Mercenary Capital” Debate

A major wave of Web3 growth (notably in DeFi) was catalyzed by token distribution mechanisms: initial coin offerings (Howell, Niessner, & Yermack, 2020), liquidity mining (issuing newly-minted tokens to users who deposit assets into a protocol, as a recruitment incentive), staking rewards (Schär, 2021; issuing tokens to participants who lock existing tokens to secure the protocol), and user rebates, designed to bootstrap two-sided markets. A key controversy is whether these incentives generate durable adoption or merely transient participation driven by extrinsic rewards.

Empirical and conceptual analyses increasingly emphasize the risk of “mercenary capital” (Schär, 2021): incentives attract liquidity or usage temporarily, but participation exits when emissions decline, leaving weak organic demand. While the quality of evidence varies across

sources, the consistent theoretical point is that token incentives can create short-run growth while undermining long-run sustainability unless paired with credible sinks, retention design, and governance discipline. This paper integrates that debate into a more general claim: **incentives are political economy**, not simply growth marketing, because they distribute resources, create winners and losers, and can entrench power.

Academic work on liquidity mining (Schär, 2021; Roughgarden, 2021) explicitly treats it as an innovative mechanism that changes investor behavior and protocol dynamics, suggesting a need for systematic evaluation rather than anecdotal interpretation.

2.4 DAO Governance: Concentration, Delegation, and the “Decentralization Illusion”

The legitimacy of tokenized governance depends not only on formal voting rules but also on power distribution and participation (Aramonte et al., 2021; Nadler & Schär, 2020) structure. Decentralized Autonomous Organizations (DAOs; on-chain organizations whose governance rules and treasury controls are encoded in smart contracts and exercised via token-weighted voting) and similar token-governed protocols face a central challenge: token-based governance can produce concentrated influence (Barbereau et al., 2023) including “whales” (very large individual holders), delegates (parties who receive voting power from many other holders), or insiders, leading to outcomes that are formally decentralized but substantively centralized. The political-theory foundations of decentralized autonomous organizations have been developed across two complementary lines of work: an early political-theory line (DuPont, 2014, “The Politics of Cryptography”; Atzori, 2017) treating DAOs as new constitutional sites, and a more recent commons-governance line (Schneider, 2024, “Governable Spaces”) that engages Ostromian polycentricity directly and provides a synthesis between DAO design and commons-management scholarship. This paper uses that literature to motivate measurement of who can actually exercise governance power; complementary work in corporate-governance theory (Tirole, 2001, on managerial accountability and minority protection) provides a structurally-parallel reference for the principal-agent concerns that recur in delegated token governance.

The Bank for International Settlements (BIS) argues (Aramonte et al., 2021) that DeFi contains a “decentralization illusion,” emphasizing that protocols marketed as decentralized frequently concentrate decision-making power among a small number of token holders, developers, or validators. This critique is significant for this paper because it links governance concentration to systemic risks and questions the legitimacy of “decentralized” claims.

Gersbach, Schneider, and Shahkar (2024) analyze vote delegation and find it can meaningfully change outcomes under highly unbalanced voting rights but has negligible impact under balanced systems, suggesting delegation is not a universal solution to legitimacy problems. Weidener, Laredo, Kumar, and Compton (2025) provide a systematic scoping review of delegated voting in DAOs, synthesizing that voting power concentration in DAOs typically exceeds token-holding concentration because delegation chains amplify the effective influence of large delegates beyond their direct token holdings, reinforcing that delegation can either mitigate or entrench concentration depending on design. Hall and Miyazaki (2024) analyze liquid democracy across 18 DAOs, finding that only 17% of tokens are delegated yet the top five delegates routinely capture over 50% of effectively voted weight, with high-SSPI delegates more likely to signal early and cascade alignment; the 1.6x to 9.9x delegation amplification ratios

documented in Section 4.5 are consistent with this clumping pattern. Esposito, Tse, and Goh (2025) evaluate quadratic voting and delegated governance against Ostrom's common-pool resource (CPR) framework, documenting that tokenization alone is insufficient for inclusive governance.

Most directly relevant, Cong, Rabetti, Wang, and Yan (2025) document that blockvoters control 76.2 percent of voting power across the DAOs they study, while proposal managers engage in insider trading yielding an average market-adjusted return of 9.5 percent. They argue that the combination of concentrated voting power and insider trading is value-destroying in the long run, particularly during crisis periods when governance becomes central to organizational survival. These stylized facts on persistent concentration in DAOs and governance token trading dynamics reinforce the need for institutional safeguards rather than assuming "token voting" equals democracy. Panel evidence from DAO-quarter data documents that effective governance control concentrates as proposal volume scales, consistent with monitoring costs outpacing broad-based participation (Tchente, 2025); the cross-sectional allocation null documented here (Pearson $r = 0.05$, $p = 0.76$, $N = 37$) is compatible with this scale-driven concentration mechanism operating independently of initial token distribution.

Calzada (2025) corroborates these patterns ethnographically across seven Web3 ecosystems, documenting whale dominance and voter turnout below 10% in major DAOs.

The governance concentration patterns documented in Section 3 are consistent with Williamson's (1999) "probity hazard": token holders can comply with smart contract rules while extracting value that contradicts the protocol's institutional mission, a risk that intensifies as governance power concentrates in fewer addresses.

Voting power inequality in Compound, Uniswap, and ENS exceeds wealth inequality in real-world economies, with voting power more unequally distributed than income in any country (Fritsch, Müller, & Wattenhofer, 2024), a pattern consistent with the delegation amplification documented in Section 4.5; notably, these concentrated delegates rarely overturn community-preferred outcomes, suggesting that concentration creates domination risk rather than routine domination. Across 10,639 proposals in 151 DAOs, three or fewer entities exert effective control over most governance decisions (Appel & Grennan, 2023), consistent with the insider dominance this paper documents at the token-holder level.

Together, these literatures (see also Feichtinger et al., 2023; Rikken et al., 2023) motivate two design imperatives advanced in later chapters: (i) governance systems require **contestability and due process** beyond raw voting, and (ii) legitimacy depends on measurable concentration constraints and multi-constituency representation.

Han, Lee, and Li (2023) formalize the governance consequences of ownership concentration, showing that whale dominance reduces platform growth but that staking mechanisms and token illiquidity can mitigate it. The finding that ownership concentration is structurally harmful, not merely descriptive, motivates the normative framework developed here: if concentration damages governance quality, the political-philosophical traditions that specify what good governance requires become directly relevant to token system design. Han, Lee, and Li (2024) provide a comprehensive review of DAO governance mechanisms and the emerging agency problem between large and small token holders.

2.4.1 Cooperatives as the Nearest Institutional Ancestor

Among established institutional forms, the cooperative bears the closest structural resemblance to the tokenized protocol. Both organize collective production without a single owner who captures surplus; both distribute governance rights to participants rather than to external shareholders; both face the challenge of sustaining voluntary participation under conditions where exit is available. The cooperative literature (Hansmann, 1996; Birchall, 2011) identifies these properties as the defining features of member-owned enterprise, distinguishing cooperatives from investor-owned firms. The International Cooperative Alliance's seven cooperative principles (ICA, 1995), including voluntary membership, democratic member control, and concern for community, map onto legitimacy constraints derived from political philosophy.

Five structural differences separate protocols from cooperatives. First, cooperatives require application and capital contribution; protocols permit permissionless entry under pseudonymity. Second, cooperative ownership is purchased; well-designed protocol ownership is earned through contribution (staking, hardware deployment, service delivery), a property Scholz (2016) identifies as the platform-cooperativism premise. The governance concentration data presented in Section 3 indicate that most protocols have not collapsed the capital-labor boundary in practice: holding HHI values span from cooperative-grade distribution (Anyone Protocol 0.013) to single-party control. Third, cooperatives enforce through bylaws, elected boards, and courts; protocols enforce through code, gaining credible neutrality at the cost of contextual judgment unless appeal mechanisms are designed in. Fourth, cooperatives operate within one jurisdiction's cooperative law; protocols operate across jurisdictions simultaneously, making Ostromian polycentricity (developed in Section 2.6) a structural necessity rather than a design preference. Fifth, cooperative fiscal discipline operates through dues and patronage dividends; protocol fiscal discipline operates through emissions and burns, with the Spend-to-Reward ratio the protocol equivalent of cooperative fiscal sustainability.

These distinctions define the contribution space: protocols inherit the cooperative's legitimacy challenge (governing collective production without a centralized principal) while adding permissionless entry, programmable enforcement, and borderless jurisdiction as additional design constraints. The polycentric framework in Section 2.6 takes up this expanded design space; the empirical results in Section 3 indicate how far current protocols remain from satisfying it.

2.5 DePIN and Tokenized Physical Infrastructure: Demand Coupling and Service Verification

DePIN expand tokenized coordination into the physical world (Bane & Gala, 2025, 2026): hardware deployment, service provision, and maintenance. DePIN intensifies the institutional nature of tokenomics because it coordinates capital formation (hardware investment), labor (operations), and service quality (SLA performance; SLA stands for service-level agreement, a contractually-specified threshold for network uptime, latency, or coverage that operators must meet to earn rewards). The result is a system that resembles public infrastructure governance more than application monetization.

Recent taxonomies formalize DePIN as a distinct coordination category with three architectural layers spanning distributed ledger, cryptoeconomic, and physical infrastructure dimensions (Ballandies et al., 2023), while Alshater (2026) synthesizes the token flywheel mechanism through which burn-and-mint equilibria (a DePIN-specific design pattern in which service users burn tokens to pay for service while operators receive newly-minted tokens as reward; the balance between burns and mints determines whether token supply contracts or expands with demand) bootstrap supply-side hardware deployment.

A key design pattern in DePIN is dual-asset economics (Voshmgir, 2020) separating speculative governance/stake tokens from stable-value usage credits. Helium provides a prominent example: network usage is paid in Data Credits, a non-transferable utility token used for service payments and onboarding actions.

This dual-asset structure motivates a general principle developed later: stable-value usage pricing reduces user friction, while governance/stake tokens enable long-term alignment and security. Separate simulation work stress-tests DePIN mechanism designs under adversarial conditions using cadCAD (a Python-based simulation framework), showing that mechanism-level choices (emission model, slashing parameters where “slashing” means automated penalty mechanisms that confiscate a validator’s staked tokens for malicious or negligent behavior) produce macro-level institutional outcomes that the normative framework developed here can evaluate.

2.6 Constitutional Political Economy and Polycentric Governance

2.6.1 Constitutions as Rule-Selection Under Constraints

Constitutional political economy (Buchanan & Tullock, 1962) treats political rules as objects of analysis and choice rather than fixed background conditions. Buchanan and Tullock frame constitutions as foundational constraints shaping later policy choices and collective decisions; this resonates with tokenized systems where “hard parameters” (e.g., inflation bounds, timelocks, emergency powers) are embedded in code. Buchanan and Tullock explicitly describe their work as analyzing the political organization of a society of free persons, using economic reasoning to examine constitutional constraints.

This literature (Buchanan & Tullock, 1962; extended in Brennan & Buchanan, 1985, on rule-following and the asymmetry between rules and the discretionary decisions they constrain; North, 1990, on the role of institutions as humanly-devised constraints structuring economic exchange) informs later chapters’ distinction between a protocol’s “constitutional layer” (hard-to-change constraints) and “policy layer” (adjustable parameters).

2.6.2 Ostrom and Polycentricity

Ostrom’s (1990) work provides one of the strongest empirical foundations for governance of commons and public service industries, building on and overturning Hardin’s (1968) “tragedy of the commons” by documenting sustained commons-management institutions that avoid both privatization and centralized state control. Her concept of polycentric governance emphasizes multiple overlapping decision centers, local experimentation, and rule diversity. Ostrom’s AER article explicitly characterizes polycentric systems as alternatives to the binary “market vs state” framing, offering an analytical framework for complex economic systems with multi-level governance. Van Vulpen and Jansen (2023) map Ostrom’s eight design principles to DAO

governance structures, providing a standardized evaluation framework for collective action sustainability. This paper draws on that literature only to motivate why governance concentration is worth measuring; optional legitimacy-scoring instruments are preserved in supplements. Li and Chen (2024) conceptualize DAOs as collective stewards of digital commons, adapting Ostrom’s design principles to emphasize equitable distribution of decision-making authority; the governance concentration data in Section 3 test whether operating protocols achieve the distributional outcomes this framing requires. De Filippi et al. (2024) analyze overlapping decision-making structures in blockchain systems through Ostrom’s polycentricity framework, identifying structure, process, and outcome dimensions that map onto the governance concentration, delegation, and participation metrics measured here.

Polycentricity (Ostrom, 1990) is particularly salient for DePIN, where service conditions differ across geography and where local knowledge is crucial for operational decisions. The paper later proposes polycentric sub-DAOs/councils as governance primitives for range-based or region-based tuning.

2.7 Synthesis and Research Gap

Across economics, computer science, political philosophy, and institutional analysis, a shared conclusion emerges: tokenized systems create constitutions and policies in code (Werner et al., 2022). Yet existing literature rarely measures who can exercise governance power after token distribution, PCA correction, and delegation. This paper addresses the gap empirically; the four-contribution structure stated in §1.1 maps to the methodology in §3 and the results in §4.

3 Data and Methods

3.1 Empirical Approach

The empirical pipeline has three parts. First, the paper constructs a 40-protocol holder-concentration cross-section and computes post-exclusion HHI, Gini, and top-holder shares. Second, it applies a protocol-controlled-address (PCA) exclusion typology before computing concentration, so staking contracts, burn addresses, treasuries, vesting contracts, bridge custody, and exchange custody do not count as independent governance holders. Third, it computes voting-HHI for protocols with sufficient delegation or voting data and compares voting concentration with post-exclusion holding concentration.

The analysis is descriptive. It tests whether observed protocol characteristics are associated with concentration in a March 2026 snapshot. It does not claim causal identification. Longitudinal panel and event-study designs are specified as future work.

Sample flow (March 2026 snapshot). The full cross-section includes $N = 40$ protocols with post-exclusion holding HHI (Table 3). The allocation battery uses $N = 37$ protocols with initial insider allocation data. The headline DePIN-vs-DeFi contrast uses 15 DePIN and 15 DeFi protocols. Subsidy analysis uses $N = 23$ protocols with non-zero subsidy under either Token Terminal or on-chain metrics, and $N = 22$ when Livepeer is excluded. Voting-HHI comparison uses $N = 13$ protocols with sufficient governance data (Table 6).

3.2 Units of Analysis and Sampling Strategy

Primary unit: **protocol-token system** (including its governance institutions). Secondary unit: **governance episodes/events** (emission changes, fee switches, unlocks, enforcement controversies, governance reforms).

Sampling principles. Protocols were selected to (i) span multiple categories (L1/L2, DeFi, infrastructure/oracles, DePIN, social), (ii) include “dead” or stagnated projects to reduce survivorship bias, and (iii) stratify by chain and launch epoch to control for regime effects.

Sample-selection criteria. Protocols were chosen to maximize variation across three dimensions: sector (15 DeFi, 15 DePIN, 8 L1/L2 infrastructure protocols where L1 denotes a base-layer blockchain like Ethereum and L2 denotes a scaling-layer chain that settles to L1, 2 failed or stagnated protocols included to avoid studying only survivors), maturity (2017-2024 launches), and governance model (token-weighted, delegate, futarchy where futarchy means governance through prediction-market betting on the outcomes of proposed policies). The expanded 40-protocol sample was constructed by adding protocols from the Dune Analytics ecosystem (Dune is a community blockchain-data platform where analysts publish SQL queries that aggregate on-chain transactions into human-readable tables) that met three criteria: (i) publicly traded governance token, (ii) holder data queryable via Dune or Helius DAS API (a Solana-blockchain data service that provides holder-balance queries via the Digital Asset Standard), and (iii) at least 50 non-zero-balance holders after protocol-controlled address exclusion. MetaDAO is included as the sole futarchy-governance protocol in the sample. Solana-native tokens (META, DRIFT, GRASS, W) required a separate data collection pipeline (Helius DAS API) because the primary Ethereum-focused indexer undercounted Solana holders by two orders of magnitude; details are in [Supplementary File S6](#).

Holder-list cutoff. The top-1,000 holder cutoff follows established convention in the token-governance concentration literature (Sai et al., 2021 use comparable cutoffs across blockchain governance studies). The cutoff captures the holder population materially relevant to governance decisions; long-tail holders below this threshold collectively hold less than five percent of supply across all sample protocols. Robustness across cutoffs is supported by the Top-1%, Top-5%, and Top-10% columns in Table 3: rank-ordering across protocols is preserved at all three cutoffs (Spearman rho > 0.95 between any two columns), indicating that absolute HHI values may shift modestly with deeper-tail inclusion but cross-protocol comparison remains robust to the cutoff choice. The canonical methodology pulls top-(1000 + number of PCAs) holders by balance and applies PCA exclusion to yield 1000 post-exclusion non-PCA holders; Table 3 N column reflects this pre-exclusion sample size (e.g., Optimism N = 1007 reflects top-1007 pre-exclusion sample with 7 PCAs excluded to yield 1000 non-PCA holders for HHI computation). For protocols with fewer than 1,000 total unique non-zero holders (e.g., Token Engineering Commons with N = 742), the actual holder count is reported as the ceiling rather than the 1,000+PCAs convention.

Gini as ledger-architecture artifact, not governance signal. Inequality (Gini) is near-universal across all sectors (range 0.52 to 0.99), while concentration (HHI) varies by an order of magnitude (0.005 to 0.199 after excluding protocol-controlled addresses). Prior work applying the Gini coefficient to eight major cryptocurrencies found inequality levels comparable to real-world economies (Sai, Buckley, & Le Gear, 2021); because permissionless address creation produces millions of near-zero-balance wallets that inflate Gini toward 1.0 regardless of

governance structure, the near-universal inequality this sample documents is a measurement artifact of ledger architecture rather than a governance signal, reinforcing the case for HHI as the primary concentration metric.

Reading Table 3. Top-1% and Top-10% columns report literal cumulative shares: Top-1% is the supply share of the single largest holder; Top-10% is the cumulative share of the top 10 holders, both computed as percentages of the post-exclusion top-1,000-holder sample total. The Top-5% column (in the canonical CSV; not displayed in this table) reports the top 5 holders cumulative share under the same convention. For protocols with fewer than 1,000 unique non-zero holders after exclusion, the same literal-top-N convention applies to the actual N. For protocols with more than 1,000 unique non-zero holders post-exclusion, the top-1,000 holders are sampled and the literal top-1 and top-10 shares within that top-1,000 sample are reported. The HHI and Top-N values in Table 3 reflect the canonical top-1,000 pre-exclusion computation; strict application of the (1,000 + PCAs) pre-exclusion convention reported in the N column would yield HHI shifts below 0.0001 and Top-N shifts below 0.1 percentage points, immaterial for the cross-protocol comparisons reported in Section 4.

3.3 Data Sources (Replicable, Multi-Disciplinary)

Four data streams inform the analysis: on-chain data (Dune Analytics, Helius DAS API), protocol documentation (whitepapers, governance forums, parameter docs), analytic aggregators (Token Terminal, DeFiLlama, CoinGecko), and public academic and policy sources. A systematic data coverage gap between DeFi and DePIN protocols (Token Terminal covers approximately 100% of DeFi but only 25% of DePIN) constrains cross-sector regression analysis. Full source provenance, exclusion methodology, the protocols Token Terminal does not cover, and the resulting partial-data subsets are documented in Supplementary File S6.

Snapshot date discipline for governance-power data: the Table 6 voting HHI values are computed from a March 2026 snapshot of Tally and Snapshot governance data, consistent with the March 2026 holder-list snapshot used for Table 3 holding HHI. Voting-power distributions can drift materially over multi-month windows: replication runs using fresh Tally and Snapshot data should expect comparable but not identical voting HHI values, and rank ordering across protocols is the more durable inferential property than absolute magnitude. Section 4.6 reports a PCA-symmetric voting-HHI sensitivity check.

Voting-HHI sampling methodology. The Table 6 voting-HHI column uses comprehensive sampling on both Tally and Snapshot sides: Tally pulls top-1,000 delegates per protocol via paginated GraphQL queries (Aave pulled at all- delegates depth $N = 150,911$ as a methodology-floor sensitivity anchor); Snapshot pulls ALL proposals in the 12-month rolling window 2025-05-22 to 2026-05-22 and aggregates unique voters across all proposals weighted by maximum voting power per voter. For high-volume protocols (ARB, COMP, LDO), this comprehensive sampling yields large unique-voter counts (ARB $N = 5,466$; LDO $N = 333$; COMP $N = 138$). Low-volume protocols (DIMO and WeatherXM with 4 or fewer proposals total) experience no methodology-induced N expansion; DIMO $N = 10$ is genuine across the protocol's entire 12-month proposal history. Solana- native protocols (JUP, HNT, DRIFT) use on-chain governance program parsing because Tally is not deployed for Solana protocols (JUP via Dune `jupiter_solana.govern_call_setvote` per-signer max-weight aggregation; HNT and DRIFT via Helius RPC parsing of SPL Governance VoteRecordV2 accounts).

The asymmetric capping reflects what each platform’s N represents: Tally caps at top-1,000 to truncate the deadweight long tail of registered-but-inactive delegates (validated as immaterial to HHI by the Aave all-delegates sensitivity anchor at N = 150,911); Snapshot and Solana on-chain N are uncapped because every entry is by construction an active participant within the window, so no analogous truncation is warranted.

Protocol	Source	N	Notes
DIMO	Snapshot	10	all 4 proposals in window; genuine sample size
Lido	Snapshot	333	all 41 proposals in window
WeatherXM	Snapshot	73	3 proposals in window
Compound	Snapshot	138	all 24 proposals in window
Aave	Tally	150,911	all-delegates pull (methodology-floor sensitivity anchor)
Uniswap	Tally	1,000	top-1,000 delegate ceiling
Arbitrum	Snapshot	5,466	all 39 proposals in window
Optimism	Tally	1,000	top-1,000 delegate ceiling
GMX	Tally	1,000	top-1,000 delegate ceiling
ENS	Tally	1,000	top-1,000 delegate ceiling
Jupiter (JUP)	Solana on-chain	128,476	per-signer max-weight aggregation on jupiter_solana.govern_call_setvote
Helium (HNT)	Solana on-chain	1,835	VSR position-state reconstruction via

Protocol	Source	N	Notes
			Helius RPC
Drift (DRIFT)	Solana on-chain	641	VoteRecordV2 parsing via Helius RPC

Table 1. Per-protocol active-voter sample size and methodology source for the Table 6 voting-HHI computation. Source-specific N reflects either Tally delegate-registry depth (top-1,000 ceiling; Aave at all-delegates anchor), Snapshot voter-pool union across the 12-month proposal window, or Solana on-chain governance program parsing.

3.4 Codebook and Variables

Insider taxonomy: the analysis distinguishes two related but distinct constructs. Insider allocation (insider_pct) is the regression covariate measured at token generation event (TGE; the moment a protocol first mints and distributes its token to a defined set of recipients), defined as team_pct + investor_pct from primary tokenomics sources. Insider count fraction in top-10 holders is a descriptive mechanism variable measured post-distribution from holder snapshots, defined as the post-exclusion fraction of wallets in the top-10 classified as insider via a three-tier procedure: Tier 1 uses Blockscout verified- contract labels and Dune exchange tags; Tier 2 uses contract bytecode matching for SafeProxy and GnosisSafe (multi-signature wallets where several authorized parties must each sign a transaction to authorize fund movement, commonly used as protocol treasuries) and vesting contracts; Tier 3 uses Etherscan (an Ethereum blockchain explorer) named-address matching against documented foundation, treasury, and team wallets.

Protocol-controlled addresses (PCAs) are excluded from holding HHI computation entirely at the top of the pipeline (133 address exclusions across 38 protocols documented in the exclusions log; 125 unique addresses with five recurring across multiple protocols) and are not counted as insiders for the purposes of the count fraction. Solana-native protocols carry an acknowledged labeling gap: io.net classifies as zero of ten insiders because Solana lacks Blockscout-style contract labels at the scale available on Ethereum and EVM (Ethereum Virtual Machine; the execution environment for Ethereum-compatible smart contracts) Layer 2 networks. This is a classification limitation rather than evidence of distributed governance and is flagged in Section 5.7.

We code each protocol across standardized dimensions via a machine-readable table capturing metadata, supply and issuance rules, distribution, value routing and sinks, staking and lockup parameters, governance architecture, decentralization metrics, incentive windows, and data/identity posture. Full field definitions and the codebook schema are documented in [Supplementary File S1](#). Optional governance-scoring instruments preserved during the broader project (Supplementary Files S2 and S3) are not used in the main-text regressions, tables, or headline findings.

3.5 Outcome Variables

Three primary metrics are used. The Herfindahl-Hirschman Index (HHI), ranging from 0 (dispersed) to 1 (monopoly), measures overall token concentration; values above 0.25 indicate high concentration by antitrust standards. Section 4.3 notes no protocol exceeds this threshold. The Gini coefficient captures distributional inequality. Top-N ratios (Top-1, Top-5, Top-10 share) complement HHI by identifying specific whale addresses.

Full metric definitions, computation procedures, and Dune query specifications are provided in [Supplementary File S1](#).

3.6 Testability and Empirical Pipeline

The measurement claims are testable in principle through cross-sectional governance analysis (this paper), longitudinal panel models, event studies around governance decisions, and qualitative process tracing. This paper reports the cross-sectional governance concentration analysis. The full quantitative pipeline, including panel regressions, event-study specifications, and robustness protocols, is documented in supplementary material ([Supplementary File S5](#)) for replication and extension by future researchers.

3.7 Ethical and Practical Considerations

All data are drawn from public sources: on-chain transactions, published governance proposals, and open protocol documentation. No private or proprietary data are used. Governance token distribution data are computed from Dune Analytics multi-chain queries. The author served as Senior Investment Analyst at Borderless Capital, a cryptocurrency-focused investment firm, through February 2026 and retains no decision-making role, carried interest, or position-dependent compensation with the firm. The author may hold small personal positions in some of the protocols analyzed; positions were established outside the research period and are not position-dependent on Borderless Capital. A complete disclosure is provided in the Conflict of Interest statement accompanying this submission.

3.8 Generalizable Protocol-Controlled-Address (PCA) Exclusion Methodology

The protocol-controlled-address (PCA) exclusion methodology applied across Sections 3.2 through 3.4 generalizes beyond this paper's specific protocol set. Any cross-protocol governance HHI analysis on ERC-20 (the Ethereum standard for fungible tokens) or SPL (Solana Program Library; Solana's analog for fungible tokens) token holders should apply the five PCA classes documented below before computing concentration metrics. The methodology contribution is not just the specific exclusions applied to this paper's 40-protocol sample ([Supplementary File S6](#) contains the address-by-address documentation), but the generalizable typology that future replication and extension work can apply uniformly.

Class 1: Burn-destination addresses. Canonical burn addresses (0x000...000 and 0x000...000dead on EVM chains; chain-specific patterns; protocol-specific burn destinations confirmed via official documentation, Blockscout/Etherscan labels, or Alchemy/Arkham labels) should be excluded from governance HHI computation. These addresses hold tokens that are inert by

construction and cannot participate in governance under any condition; including them in the holding distribution inflates the holding denominator and understates effective concentration. The Uniswap 0x000...000dead exclusion (102.46M UNI, 11.27 percent of supply at measurement) is the canonical case in this paper; the Hyperliquid burn-pattern addresses (0x222...222 and 0xfefe...fefe) were caught by a confirmed-pattern methodology. The universal burn-rule applied here is generalizable to any ERC-20 or SPL token governance analysis.

Class 2: Foundation and treasury custody. Protocol-controlled foundation and treasury custody addresses should be excluded. These addresses hold tokens that are governance-relevant in principle (the foundation may delegate or vote with them) but are structurally distinct from independent holder concentration: the foundation position represents institutional rather than community ownership. Specific cases in this paper: Optimism Foundation GnosisSafe (0x2a82ae142...; 32 percent of pre-exclusion top-1000 OP); Arbitrum FixedDelegateErc20Wallet (0xf3fc178157...; 30 percent of pre-exclusion top-1000 ARB); WeatherXM treasury (74.8 percent of pre-exclusion top-1000 WXM); DIMO Foundation (52.3 percent of pre-exclusion top-1000 DIMO). The methodology requires per-protocol identification typically via Blockscout verified-contract labels, GnosisSafeProxyFactory creation records, or foundation disclosures.

Class 3: Staking-aggregation contracts. Staking contracts that aggregate underlying staker holdings into a single custody address should be excluded from holding HHI computation, with the caveat that underlying stakers retain pass-through voting power. The stkAAVE staking contract (0x4da27a545c0c...; 21.9 percent of post-PCA-exclusion top-1000 AAVE) is the canonical case; AAVE stakers vote through stkAAVE rather than the contract itself making governance decisions. The sGMX RewardTracker plays an analogous role for GMX. The exclusion is methodologically warranted at the holding-HHI computation step because the contract is not an independent governance participant; the exclusion at the voting-HHI computation step (Section 4.6 PCA-symmetric robustness check) is methodologically distinct: stkAAVE votes as a single delegate per Tally's data model, even though the underlying voters are dispersed. The paper treats the holding-side exclusion as canonical and reports the voting-side exclusion as a sensitivity check.

Class 4: Bridge custody and migration addresses. Bridge custody addresses (Wormhole token bridge for cross-chain wrapped tokens; the GRT BridgeEscrow for L1-to-Arbitrum custody; LayerZero canonical bridges) hold tokens in transit or as cross-chain accounting positions; these are not governance-participating holders. Migration contracts (the SYRUP Migrator for MPL-to-SYRUP migration; the Polygon StakeManagerProxy for MATIC-to-POL transition) hold tokens during migration windows and should be excluded.

Class 5: Centralized exchange custody. Centralized exchange (CEX) hot and cold wallets that custody customer deposits should be excluded from governance HHI computation. Customer-deposited tokens held in CEX custody are not governance-participating holdings: the exchange holds custody of tokens belonging to many individual customers, and exchange operational practice is to abstain from governance voting on customer-held tokens. Specific cases in this paper: Binance hot wallets across OP, ENS, GRT, AAVE, COMP, UNI, ARB, and other protocols (the canonical Binance address 0xf977...41acec is the most common case); Bithumb 162 (0x54d1...22cbf9) and Upbit 59 (0x377b...190873) on AXL; analogous patterns on other protocols. CEX address identification relies on public name tags (Etherscan, Arkham, hildobby

compilations) and behavioral signals (high transfer volume, many counterparties, no governance participation history).

Across the five classes, this paper applies 133 PCA exclusions across 38 protocols (Supplementary File S6 and the supplementary exclusions log dataset). The generalizable methodological contribution: any future cross-protocol governance HHI analysis should apply the five-class PCA exclusion uniformly before reporting holding concentration metrics. The specific addresses depend on the protocol sample; the methodology is sample-independent.

4 Results

4.1 Scope and Claims Boundary

The cross-section covers 40 protocols across four sectors (DePIN, DeFi, infrastructure, social). Table 2 delimits the claims boundary: all findings are descriptive within the observed scope. The governance data represent a March 2026 snapshot, not a time series. Causal claims about the relationship between institutional design and governance outcomes require the longitudinal evidence beyond the scope of this paper. All governance token distributions were queried from Dune Analytics multi-chain queries. Concentration metrics are point-in-time snapshots; trajectory analysis requires the longitudinal extension specified in the Future Research discussion (Section 5.8).

Claim Evidence Grade Pointer

PCA exclusion corrects systematic HHI inflation Methodological Section 3.8; S6

Allocation null with insider-retention contrast Descriptive Section 4.4; Table 4

DePIN concentration exceeds DeFi after correction Descriptive Section 4.3; Table 5


Delegation often amplifies voting above holdings Descriptive Section 4.5; Table 6

Subsidy association is Livepeer-driven only Descriptive (null) Section 4.4; S8

Table 2. *Claims boundary: what this paper does and does not claim.*

4.2 Cross-Protocol Concentration

Governance concentration spans two orders of magnitude across the 40-protocol sample (Table 3): HHI ranges from 0.005 (Hyperliquid) to 0.199 (Livepeer), with a median near 0.04. DePIN protocols cluster in the upper half of this distribution.



Source: Dune Analytics + Helius DAS API (March 2026). 69 protocol-controlled addresses excluded across 20 protocols.

Figure 1. Holding-based HHI across 40 governance tokens (March 2026 snapshot; top-1,000 holders; 133 protocol-controlled addresses excluded across 38 protocols per Section 3.8 five-class typology). Bars colored by category (DeFi orange; DePIN blue; L1/L2 Infrastructure green; Social grey); sorted by category then HHI ascending. Dashed vertical line at 0.25 marks the antitrust concentration threshold. Highest HHIs cluster in DePIN (Livepeer 0.199; WeatherXM 0.148; io.net 0.125; Aethir 0.095); DeFi and L1/L2 protocols sit predominantly below 0.05.

A measurement caveat applies to interpretation: the 0.25 HHI antitrust threshold is borrowed from market-concentration analysis (DOJ Horizontal Merger Guidelines) and is reported here as a benchmark for descriptive interpretation only. A market-concentration threshold is not

equivalent to a democratic-legitimacy threshold; a protocol may have HHI below 0.25 while remaining politically dominated through quorum rules (the minimum fraction of voting power required for a proposal to pass; a high quorum may concentrate decisive influence in whoever holds enough tokens to meet it), delegation patterns, or proposal-rights concentration. The HHI captures one dimension of concentration (post-exclusion holding distribution) and is not a sufficient condition for distributed governance.

Livepeer is the clearest structural outlier; excluding it leaves a distribution in which DePIN still exhibits higher concentration than DeFi but without a dominant leverage point.

Concentration metrics are computed from Dune Analytics multi-chain queries and the Helius DAS API for Solana SPL tokens (March 2026 snapshot, top-1000 holders per protocol), after applying 133 protocol-controlled address exclusions across 38 protocols (staking contracts, exchange custodians, vesting locks, protocol treasuries; see [Supplementary File S6](#)). Solana-native tokens use getTokenAccounts with page-based pagination to correct systematic undercounting in Dune’s indexed tables. Replication scripts and input CSVs are listed in Supplementary File S5.

Three additional DePIN protocols are included to improve statistical power for the subsidy analysis: Aethir (ATH, Ethereum, GPU compute, \$156M annualized revenue (Token Terminal, April 2026)), Hivemapper (HONEY, Solana, mapping), and io.net (IO, Solana, GPU compute, \$21M revenue). Governance concentration was computed using the same methodology (Dune for ERC-20 tokens, Helius for Solana SPL tokens, top-1000 holders, exchange-labeled and protocol-controlled addresses excluded).

Protocol	Token	Category	HHI	Gini	Top-1	Top-10	N
Hivemapper	HONEY	DePIN	0.018	0.91	5.7%	31.1%	1000
Morpheus AI	MOR	DePIN	0.0464	0.88	12.6%	55.4%	1003
Render	RENDER	DePIN	0.0268	0.88	9.9%	43.1%	1001
Grass	GRASS	DePIN	0.0352	0.76	16.0%	38.2%	1000
DIMO	DIMO	DePIN	0.0249	0.85	5.9%	39.3%	1002
Anyone Protocol	ANYONE	DePIN	0.0128	0.71	4.7%	30.1%	1003
Filecoin	FIL	DePIN	0.0221	0.74	9.3%	35.6%	1001
Helium	HNT	DePIN	0.0745	0.98	16.5%	76.4%	1001
IoTeX	IOTX	DePIN	0.0810	0.96	12.3%	80.9%	1008
io.net	IO	DePIN	0.1251	0.94	33.5%	61.2%	1000

Protocol	Token	Category	HHI	Gini	Top-1	Top-10	N
GEODNET	GEOD	DePIN	0.0605	0.91	16.2%	62.3%	1004
WeatherXM	WXM	DePIN	0.1479	0.89	35.0%	77.9%	1001
Pokt Network	POKT	DePIN	0.0899	0.94	25.7%	66.7%	1001
Aethir	ATH	DePIN	0.0948	0.94	19.7%	67.6%	1008
Livepeer	LPT	DePIN	0.1989	0.94	42.9%	73.7%	1002
Hyperliquid	HYPE	DeFi	0.0052	0.65	2.4%	15.5%	1000
Lido	LDO	DeFi	0.0077	0.77	3.1%	18.6%	1009
Maple Finance	MPL_SYRUP	DeFi	0.0242	0.88	8.5%	40.6%	1002
Compound	COMP	DeFi	0.0092	0.84	3.5%	20.8%	1005
Balancer	BAL	DeFi	0.0290	0.87	13.2%	35.2%	1011
Rocket Pool	RPL	DeFi	0.0392	0.87	11.5%	52.7%	1002
MakerDAO	MKR	DeFi	0.0404	0.88	13.9%	50.6%	1002
MetaDAO †	META	DeFi	0.0152	0.88	6.2%	28.6%	999
GMX	GMX	DeFi	0.0647	0.92	20.1%	59.3%	1004
Aave	AAVE	DeFi	0.0128	0.83	4.3%	27.7%	1003
Ether.Fi	ETHFI	DeFi	0.0417	0.93	15.7%	47.8%	1002
Uniswap	UNI	DeFi	0.010	0.89	4.2%	22.9%	1002
Drift	DRIFT	DeFi	0.0529	0.92	19.3%	49.1%	1001
Jupiter	JUP	DeFi	0.0957	0.98	26.0%	69.3%	1001
Curve	CRV	DeFi	0.0144	0.79	6.6%	31.2%	1010
Axelar	AXL	L1/L2/Infra	0.0268	0.85	9.0%	41.7%	1004
LayerZero	ZRO	L1/L2/Infra	0.0143	0.89	4.0%	29.9%	1002
Wormhole	W	L1/L2/Infra	0.0115	0.86	7.0%	19.0%	1001

Protocol	Token	Category	HHI	Gini	Top-1	Top-10	N
Polygon	POL	L1/L2/Infra	0.0348	0.92	12.1%	49.5%	1002
Arbitrum	ARB	L1/L2/Infra	0.0119	0.73	8.6%	22.4%	1007
The Graph	GRT	L1/L2/Infra	0.033	0.92	13.1%	40.2%	1004
Optimism	OP	L1/L2/Infra	0.0093	0.77	5.7%	21.0%	1007
ENS	ENS	L1/L2/Infra	0.0494	0.86	16.8%	45.2%	1004
Token Engineering Commons ‡	TEC	Social	0.0278	0.91	9.2%	44.2%	742
Gitcoin ‡	GTC	Social	0.0215	0.89	7.1%	36.8%	996

Table 3. Governance concentration cross-section (March 2026 snapshot; Dune Analytics + Helius DAS API; top-(1,000 + number of PCAs) holders per token pre-exclusion, yielding 1,000 post-exclusion non-PCA holders for HHI computation; 133 protocol-controlled-address exclusions across 38 protocols per Section 3.8 five-class typology; N column reflects pre-exclusion sample size). Sorted by category then HHI ascending. HHI values measure raw token-holding concentration, not voting power; vote-escrowed concentration (Curve veCRV, Balancer veBAL) is expected to be higher due to lock-duration multipliers (Section 4.5.2). Hyperliquid is excluded from the subsidy regression because 97% of trading fees route to an Assistance Fund that permanently burns HYPE tokens (sink-mechanism in framework). Aethir HHI (0.095) reflects May 2026 Dune top-1008 pull with eight PCA exclusions; Hivemapper + io.net rows from Helius DAS API.

Table 3 notes:

† = futarchy governance (MetaDAO).

‡ (in protocol name column) = survivorship-bias control (TEC, Gitcoin).

Methodology notes: HHI for Curve reflects raw CRV holdings; veCRV-weighted voting concentration is approximately 0.26 computed from cumulative CRV deposit volume across the top-50 historical veCRV lockers (Convex Finance, an aggregator that pools user-deposited CRV into a single locked veCRV position and votes as a meta-governance layer on behalf of depositors, holds 48% of the top-50 deposit-weighted share; Yearn Finance, a yield-aggregation protocol that runs a similar meta-governance position, holds 14%; Stake DAO holds 9%; the remaining top-50 lockers each hold less than 4%); the methodology is discussed in Section 4.5.2. HHI for Uniswap reflects post-burn-rule exclusion of 0x000...dead (102.46M UNI, 11.27% of supply). Top-1% and Top-10% columns for AAVE, UNI, ARB, GRT, and OP report post-exclusion shares, applying the same protocol-controlled-address exclusion methodology used for HHI computation.

HHI ranges from 0.005 to 0.199 across the full sample. Among DeFi governance tokens, HHI ranges from 0.005 (Hyperliquid, after excluding the Assistance Fund and burn addresses) to 0.096 (Jupiter). Curve's raw CRV holding HHI of 0.014 reflects the post-Convex-exclusion distribution; the veCRV-weighted voting concentration is substantially higher (approximately 0.26 computed from cumulative deposit volume across the top-50 veCRV lockers; Section 4.5.2), and is discussed in Section 4.5.2. DePIN protocols span a similar range: 0.018 (Hivemapper) to 0.199 (Livepeer); Morpheus AI registers 0.046 after PCA exclusion (per Table 3 row 0.0464).

Several protocols exhibit much lower concentration in the post-exclusion top-1,000 sample than in the raw holder list. Applying the holder-list correction for Solana tokens (Helius DAS API enumeration versus the systematically-undercounting Dune indexer) and the PCA exclusion methodology (38 protocols with documented exclusions; Section 2.10.10), the DePIN-DeFi concentration gap is computed on a holder sample materially different from the naive top-1,000 alternative. For DRIFT, the raw holder distribution yields HHI 0.419 against the post-exclusion value of 0.053 (Helius enumerates 27,888 distinct holders versus 104 retrieved from the Dune indexer). Wormhole (W) moves from 0.126 (raw) to 0.012 (post-exclusion); GRASS moves from 0.132 to 0.035. The exclusion methodology ([Supplementary File S6](#)) is particularly consequential for protocols with large foundation treasuries or exchange custodian addresses in the top-10 holders.

Several protocols show that low governance concentration is achievable across sectors. Hyperliquid (HHI 0.005), Lido (0.008), Compound (0.009), and Uniswap (0.010) in DeFi; Optimism (0.009), Arbitrum (0.012), Wormhole (0.012), and LayerZero (0.014) in L1/L2 infrastructure; and Hivemapper (0.018) in DePIN all achieve HHI below 0.02. These counterexamples indicate that low governance concentration is achievable across multiple sectors, constraining explanations that treat high concentration as structurally inevitable in any single sector. The fourth sector, social token protocols, is too small in the sample ($N = 2$) to support independent sector-level claims.


Three measurement choices affect these results and merit attention as scope conditions. First, Top-N cutoffs at 1,000 holders may undercount long-tail holders; increasing the cutoff to 10,000 addresses would change absolute HHI values but is unlikely to change the rank ordering. Second, the snapshot methodology captures a single point in time; several protocols (notably DIMO) have announced token distribution programs that may reduce concentration over subsequent months. Third, the analysis uses raw token holdings; for protocols with delegation (Compound, Uniswap), delegated voting power may differ from raw holdings. We use delegated power where available, but not all protocols instrument delegation.

4.3 DePIN vs. DeFi Contrast

DePIN protocols have approximately two and a third times the governance concentration of DeFi protocols (mean HHI 0.071 vs. 0.031, a ratio of 2.3; Mann-Whitney $p = 0.020$, Cohen's $d = 0.94$, $N = 15/15$). The gap holds at the conventional 0.05 threshold in Models 1 and 2 (Table 5) after adjusting for protocol age and valuation, and falls just above the threshold in Model 3 ($p = 0.078$) after adding the initial insider allocation control. Prior review work documents that DePIN sectors vary in tokenomic architecture (Alshater, 2026): storage networks prioritize staking, wireless coverage emphasizes service verification, compute networks focus on billing.

The concentration gap between DePIN and DeFi is not implied by this sectoral variation; it is a separate finding about who holds governance tokens.

The sector difference is robust to sample composition: the Mann-Whitney result holds across all 30 leave-one-out (LOO) iterations (p range across iterations: 0.008 to 0.034; Cohen's d range: 0.84 to 1.10; per-iteration LOO values in Supplementary File S5), and a two-sided permutation test (100,000 random reassignments of sector labels) yields $p = 0.012$. A non-parametric bootstrap on the sector mean difference (a resampling technique that estimates statistical variability by drawing 10,000 random samples with replacement from the observed data) yields a 95% percentile interval of [0.012, 0.069] HHI points (Cohen's d 95% confidence interval (CI): [0.32, 1.68]), strictly excluding zero, evidence consistent with the sector contrast being robust to resampling. One candidate mechanism: in capital-intensive networks, scale economies produce concentrated market shares (Arnosti & Weinberg, 2022). Fleet operators deploying hundreds of devices achieve lower per-unit costs than solo participants, consistent with the concentration gap observed here.



Source: Author calculations. Diamond = mean. Horizontal line = median. Individual protocols overlaid.

Figure 2. Governance concentration (HHI) by sector. DePIN mean (0.071) exceeds DeFi mean (0.031), with distributions overlapping substantially: 8 of 15 DePIN protocols fall within the DeFi range. The Mann-Whitney test ($p = 0.020$, $d = 0.94$) is robust to sample composition (significant in all 30 leave-one-out iterations). Reading the figure: box plots show the concentration distribution within each sector; the DePIN-DeFi gap is visible in median position despite overlap in tails.

DePIN governance concentration is also significantly higher than L1/L2 infrastructure protocols (mean HHI 0.024, $N = 8$): Welch t-test (a t-test variant that does not assume equal variances across the two groups) $p = 0.006$, Cohen's $d = 1.04$; Mann-Whitney $p = 0.016$. This divergence

reflects post-exclusion infrastructure HHIs: ARB (0.012), OP (0.009), and GRT (0.033) show distributed holdings once foundation and bridge addresses are excluded.

Revenue dynamics reinforce this structural contrast: during 2025's market downturn, Helium's onchain revenue grew over 800% while leading DeFi protocol fees declined 27-70% over the same period (Bane & Gala, 2026), suggesting that physical infrastructure revenue streams exhibit countercyclical properties absent in purely financial protocols.

This concentration gap is not explained by protocol maturity: Compound (2018 launch, HHI 0.009) and MakerDAO (2017, HHI 0.040) are older than DIMO (2022, HHI 0.025) and WeatherXM (2024, HHI 0.148), yet maintain far lower concentration. One candidate explanation is distribution design (Table 3): DePIN token allocations frequently reserve large shares for founding teams and early investors, while DeFi protocols have had longer periods of community distribution through liquidity mining, delegation, and governance participation incentives. However, the regression analysis below finds no significant correlation between insider allocation share and HHI (Section 4.4), suggesting that the distribution design channel operates through qualitative mechanisms not captured by the allocation percentage alone. The Anyone Protocol counterexample supports this interpretation: its low concentration (HHI 0.013) is consistent with governance mechanism choices (delegation incentives, activity-based caps, distribution schedule structure) that most DePIN protocols have not adopted.

This governance mechanism interpretation aligns with practitioner evidence from points-based token distribution programs (Sawinyh, 2025). Mature DeFi protocols distributed tokens through multi-season liquidity mining programs, safety module staking, and delegation incentives that implement sub-linear scaling (diminishing returns per marginal dollar deposited), activity-based caps, and time-weighted multipliers; these mechanisms structurally compress holding concentration over years. DePIN protocols that allocated tokens through conventional venture rounds, reserving 30 to 50 percent for founding teams and early investors with cliff-based vesting, bypassed these distribution-broadening mechanisms entirely. The governance concentration patterns in Table 3 may reflect distribution mechanism design operating through the sector channel: Table 5 shows DePIN membership is positively associated with concentration at $p < 0.05$ after adjusting for protocol age, valuation, and insider allocation. However, the regression null on insider allocation (Section 4.4) cautions against treating initial token allocation as a complete measure of distribution design: compression mechanisms (sub-linear scaling, activity-based caps, delegation incentives) operate through ongoing protocol governance, not through the initial percentage split alone.

4.4 Allocation Design and Post-Distribution Dynamics

Initial token allocation design is not associated with governance concentration. The percentage of total supply allocated to team and investors at launch shows no significant association with post-distribution HHI (Pearson $r = 0.05$, $p = 0.76$, $N = 37$; Spearman $\rho = 0.07$, $p = 0.69$). This null persists across alternative specifications: token unlock progress, float-adjusted concentration, and the interaction of allocation with protocol maturity all produce coefficients indistinguishable from zero (Table 5, Models 1 through 3; Table 4 summarizes the bivariate battery). Statistically Validated Network analysis across DeFi protocols documents that governance concentration shifts dynamically in correlation with total value locked and secondary

market valuations (Eisermann et al., 2025), consistent with the post-distribution market dynamics that render initial allocation design uninformative in the cross-section.

Methodological note on AAVE: the stkAAVE staking contract is excluded from the holding Herfindahl-Hirschman Index per the protocol-controlled-address rule, but stakers retain pass-through voting power. The reported AAVE holding HHI of 0.013 therefore understates effective governance concentration; reconstructing the staker distribution would require parsing all stkAAVE deposit events, which is deferred to follow-up work. The companion Tally-sourced AAVE delegation HHI of 0.058 in Table 6 partially closes this gap by capturing concentration in delegated voting rather than raw token holdings.

Longitudinal panel evidence consistent with the cross-sectional null. A companion quarterly HHI panel for 14 governance tokens over 8 quarters post-token-generation event (Supplementary File S7) shows that 11 of 14 protocols exhibit monotonic governance HHI decay over the 24-month observation window, with the three exceptions reflecting mechanism-specific lock-in: CRV (vote-escrow), COMP (reservoir-drip emissions), and ENS (claim-window distribution). The longitudinal pattern reinforces the cross-sectional finding that initial allocation is not associated with steady-state concentration: distributions tend to deconcentrate over time across heterogeneous allocation designs, indicating that the relevant institutional-design margin is post-distribution architecture (delegation programs, lock-in mechanisms, distribution-channel diversity) rather than launch-time allocation. The panel design supports the cross-section’s descriptive claim without resolving causal identification (which requires the panel-and-event-study extension specified in Section 5.8).

Covariate	Test	Statistic	p	N	Status
Insider allocation (%)	Pearson r	0.05	0.76	37	Null
Team allocation (%)	Pearson r	-0.08	0.62	37	Null
Investor allocation (%)	Pearson r	0.14	0.39	37	Null
Protocol maturity (years)	Pearson r	0.02	0.92	38	Null
Circ/total supply ratio	Pearson r	0.03	0.86	31	Null
MCap/FDV ratio	Pearson r	0.21	0.25	33	Null
Exclusion-adjusted float	Pearson r	-0.003	0.99	30	Null
Insider count fraction (top-10)	Spearman ρ	0.48	0.003	37	Significant*

Covariate	Test	Statistic	p	N	Status
Non-insider HHI, V1 (count)	Spearman ρ	0.54	0.001	34	Tautology-checked
Non-insider HHI, V2 (balance)	Spearman ρ	0.33	0.060	34	Does not survive
Subsidy ratio (cross-sector)	Pearson r	0.59	0.003	23	Fragile (LPT-driven)
Subsidy ratio (TT expanded)	Pearson r	0.12	0.61	20	Null (robustness check)
Subsidy ratio (excl. Livepeer)	Pearson r	0.07	0.77	22	Null
DePIN sector indicator	r Mann-Whitney	U=172	0.016	30	Robust (30/30 LOO)
HHI-Gini	Pearson r	0.59	<0.001	40	Significant


Table 4. Bivariate associations between protocol characteristics and governance concentration (HHI). Allocation, maturity, and float specifications are null. Insider wallet presence is significant but the balance-based measure does not survive a tautology check (Section 4.4). Subsidy is significant in levels but driven entirely by Livepeer (Section 4.6). * LOO-robust (37/37). MCap = market capitalization; FDV = fully diluted valuation.

To understand why allocation is uninformative, we classify each protocol’s top-10 post-exclusion holders using a three-tier methodology: Tier 1 identifies addresses via Blockscout labels and Dune exchange tags; Tier 2 inspects contract bytecode (SafeProxy, GnosisSafe, vesting contracts) via Blockscout’s implementation field; Tier 3 applies Etherscan named-address matching. Of 390 addresses classified across 39 tokens, 119 (30.5%) are identifiable as insider wallets (team, foundation, or investor multisigs), 67 (17.2%) are exchange addresses, 78 (20.0%) are protocol contracts, and 126 (32.3%) remain unclassified. The 78 protocol contracts identified here differ from the 133 protocol-controlled address exclusions applied to HHI computation. The 78 figure counts protocol-contract addresses across the top-10 holders of 39 tokens. The 133 figure counts excluded address inclusions across the top-1,000 holders of the 38 protocols with documented PCA exclusions (125 unique addresses; five addresses recur across multiple protocols). The samples and windows differ; the two numbers are not interchangeable.

Protocols whose insider wallets retain top-holder positions tend to exhibit higher governance concentration (Spearman $\rho = 0.48$, $p = 0.003$, $N = 37$), robust across all 37 leave-one-out iterations (ρ range: 0.45 to 0.55). However, this association has a mechanical component: insider wallets with large balances contribute directly to HHI through the squared-share calculation. To verify that the relationship is not a tautological artifact, we recomputed concentration after excluding insider addresses and renormalizing remaining shares. The count-based measure survives: non-insider HHI remains significantly correlated with the fraction of

top-10 holders that are insiders (Spearman rho = 0.54, $p = 0.001$, $N = 34$; N drops from 37 to 34 because the three zero-insider protocols, Hyperliquid, Optimism, and Arbitrum, have no insiders to exclude, so the tautology check is undefined for them), indicating that insider-heavy protocols have concentrated governance ecosystems, not merely concentrated insider positions. The balance-based measure loses significance (rho = 0.33, $p = 0.060$), consistent with it capturing the mechanical contribution of large insider stakes rather than a structural relationship. The informative finding is not the insider correlation per se but its contrast with allocation design: what matters is not how tokens were designed to be distributed, but whether the recipients retained or sold their positions.

The divergence between allocation design (null) and current insider presence (descriptive mechanism) carries a design implication. Governance concentration is not fixed at launch; it emerges from post-distribution dynamics. The allocation null (Section 4.4) is also consistent with Michels' (1911) observation that organizational complexity inevitably concentrates power in a specialized leadership class: early protocols' commitment to governance minimization, the assumption that minimal on-chain rules would produce emergent decentralization is inconsistent with the concentration patterns the ideology sought to prevent. Protocols whose insiders retain their positions exhibit more concentrated governance outcomes, regardless of how generously the initial allocation was designed. The count-based tautology check (rho = 0.54 on non-insider HHI) suggests this is not merely arithmetic: insider-heavy protocols attract or retain concentrated non-insider holders as well, consistent with a pattern where concentrated insider presence attracts or coexists with concentrated non-insider holders, rather than insider balances alone inflating the metric.



Source: Author calculations from Dune Analytics and Token Terminal (March 2026).

Figure 3. Insider allocation (%) vs. post-exclusion holding HHI for 37 regression-ready protocols. The null relationship ($r = 0.05$, $p = 0.76$) is the paper's lead finding: initial allocation design is not associated with post-distribution governance concentration. Reading the figure: the absence of a downward slope is the finding; if insider allocation drove concentration we would expect a positive slope, and if exclusion overcorrected we would expect a negative slope.

Hyperliquid provides the strongest available test of the allocation design hypothesis. The protocol accepted zero venture capital, distributed 31% of total supply via a single community airdrop (the largest in crypto history by dollar value), and operates a deflationary model where 97% of trading fees purchase and burn the native token; it has no ongoing emissions. Five protocol-controlled addresses were excluded from the HHI computation: the Assistance Fund system address (confirmed in Hyperliquid documentation), the burn address, the community rewards treasury (241 million tokens, matching the published allocation), the Foundation treasury (60 million tokens, exactly 6% of total supply), and the staking manager contract (confirmed via Hypurrsan). The resulting HHI of 0.005 is the lowest in the DeFi subsample (Lido at 0.008 is the second-lowest after the post-PCA-audit exclusions, and Compound at 0.009 the third-lowest). Despite the most community-favorable distribution architecture in the sample, this single observation cannot support a causal identification claim on its own. The cross-sectional null (Pearson $r = 0.05$, $p = 0.76$, $N = 37$) shows that protocols across the full range of allocation designs achieve comparable concentration levels, consistent with the finding that allocation design is not associated with post-distribution concentration.

The governance concentration patterns documented across the sample (Section 4.2, Table 3; HHI range 0.005 to 0.199) are descriptive, but they matter for governance audits because protocols often cite token distribution as evidence of decentralization. The corrected data show that current holder composition, insider retention, and delegation mechanics are more informative audit targets than launch allocation percentages alone.

The null pattern extends beyond allocation to financial sustainability more broadly. Three protocols with similar subsidy ratios (Aethir at 0.36x, Aave at 0.37x, io.net at 0.40x) exhibit governance HHI values of 0.095, 0.013, and 0.125 respectively. All three are revenue-dominant; their concentration spans much of the 40-protocol sample range. Financial sustainability is neither necessary nor sufficient for governance decentralization.

Table 5 reports multivariate OLS regressions of HHI on protocol characteristics. The DePIN coefficient is positive and statistically significant in Models 1 and 2 (Model 1: 0.045, $p = 0.012$, $N = 38$; Model 2: 0.053, $p = 0.036$, $N = 36$) and just above the conventional 0.05 threshold in Model 3 (0.043, $p = 0.078$, $N = 35$). No other covariate achieves $p < 0.10$. In particular, the initial insider allocation percentage used as a Table 5 covariate is not significant, consistent with the bivariate null in Table 4. This is distinct from the current insider-retention measure reported in Section 4.4 (Spearman $\rho = 0.48$), which counts how many of a protocol's top-10 holders are insiders rather than measuring the share of supply allocated to insiders at launch. The adjusted R-squared ranges from 0.15 to 0.19, indicating that sector membership accounts for more cross-sectional variation than protocol age, valuation, or allocation structure. The sample ($N = 35$ to 38 depending on specification) remains too small to reliably isolate independent effects; the pipeline specification ([Supplementary File S5](#)) documents the infrastructure for replication at larger sample sizes.

	Model 1	Model 2	Model 3
DePIN (vs DeFi)	0.045** (0.020)	0.043** (0.025)	0.043* (0.029)
L1/L2/Infra (vs DeFi)	-0.005 (0.011)	-0.002 (0.014)	-0.004 (0.016)
Social (vs DeFi)	-0.010 (0.079)	-0.020 (0.064)	-0.019 (0.146)
Protocol age (years)		0.002 (0.005)	0.001 (0.006)
Log FDV		-0.003 (0.004)	-0.003 (0.005)
Initial insider allocation (%)			0.000 (0.001)
Constant	0.032*** (0.007)	0.076 (0.089)	0.063 (0.099)
N	38	36	35
Adjusted R ²	0.173	0.190	0.151

Table 5. OLS regression of HHI on protocol characteristics. Model 1 includes sector dummies only (DePIN, L1/L2/Infra, Social; DeFi omitted). Model 2 adds protocol age (years since token genesis) and log FDV (log fully diluted valuation). Model 3 adds the initial insider allocation percentage (team plus investor share at launch). Successive models lose observations as covariate availability declines ($N = 38$ to 36 to 35). Robust standard errors (HC3; a heteroskedasticity-consistent estimator that produces valid standard-error estimates even when error variance is not constant across observations, the typical condition in cross-sectional cross-protocol data) in parentheses. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

A bootstrap robustness check supports the sector coefficient is reliably positive across resampling under the full Model 3 specification. With 5,000 bootstrap resamples of the $N = 35$ sample, the DePIN coefficient bootstrap distribution centers at +0.054 with a 95% percentile interval of [+0.001, +0.091], strictly above zero. The bootstrap-positive fraction is 98.0%. The Model 3 attenuation relative to Models 1 and 2 is attributable to partial collinearity between the DePIN sector indicator and the insider-allocation control: in the cross-section, DePIN protocols

cluster on higher insider allocation than DeFi protocols, so a fraction of the sector coefficient's signal is absorbed by the insider variable when both appear in the same model. This is consistent with the cross-sectional allocation null (Table 4, Pearson $r = 0.05$ for insider allocation alone) rather than indicating that the sector effect is fragile to covariate inclusion. The $N = 35$ multivariate sample also limits the precision with which the sector coefficient can be statistically isolated from correlated controls; the bootstrap robustness check above is the more informative robustness statement under small-sample multivariate conditions.

4.5 Voting Power vs Holding Concentration

Delegation usually amplifies governance concentration above holding levels in the protocols with comparable voting data. Table 6 compares post-exclusion holding HHI with voting-HHI for thirteen protocols. Nine protocols amplify, with ratios from 1.6x (DRIFT) to 9.9x (Compound) and a mean of approximately 4.9x. Four protocols disperse voting power below the holding baseline: ENS (0.45x), GMX (0.87x), HNT (0.35x to 0.53x), and JUP (0.057x). The holding-based HHI reported in Table 3 measures token distribution, not governance power; delegated voting data are therefore necessary for protocols where votes flow through delegates, governance platforms, or lockup-weighted voting systems. To quantify this gap, the analysis uses Tally, Snapshot, and Solana on-chain governance data from the 12 months preceding the voting snapshot.

Amplification varies by sector composition rather than by sector membership. Within DeFi, the range spans 1.6x (DRIFT) to 9.9x (Compound); DePIN spans 3.8x (WeatherXM) to 9.1x (DIMO); infrastructure spans 3.1x (Arbitrum) to 3.6x (Optimism). The four dispersion exceptions are discussed individually in §4.5.1 (ENS), §4.5.1.1 (GMX), and §4.5.4 (HNT and JUP under within-Solana heterogeneity). Sector membership is not associated with how strongly delegation concentrates power; delegate-program design is the more plausible measurement surface. DIMO's voting HHI (0.228) is 9.1 times its holding HHI (0.025), though this ratio is based on only 10 Snapshot voters and should be interpreted with caution, and WeatherXM's (0.556) is 3.8 times its holding HHI (0.148). A small number of delegates, likely foundation-affiliated or early-investor wallets, control governance outcomes despite relatively distributed token holdings.

Protocol	Category	Holding HHI	Voting HHI	Ratio	Source	N
DIMO	DePIN	0.025	0.228	9.1x	Snapshot	10
Lido	DeFi	0.008	0.050	6.3x	Snapshot	333
WeatherXM	DePIN	0.148	0.556	3.8x	Snapshot	73
Compound	DeFi	0.009	0.089	9.9x	Snapshot†	138
Aave	DeFi	0.013	0.058	4.4x	Tally§ 1	50,911
Uniswap	DeFi	0.010	0.027	2.7x	Tally	1000

Protocol	Category	Holding HHI	Voting HHI	Ratio	Source	N
Arbitrum	Infra	0.012	0.038	3.1x	Snapshot†	5,466
Optimism	Infra	0.009	0.033	3.6x	Tally	1000
DRIFT	DeFi	0.053	0.083	1.6x	Solana VSR¶	641
GMX	DeFi	0.065	0.057	0.87x	Tally‡	1000
HNT	DePIN	0.075	0.026-0.039	0.35x-0.53x	Solana VSR¶	1,835
ENS	Infra	0.049	0.022	0.45x	Tally‡	1000
JUP	DeFi	0.096	0.0055	0.057x	Solana delegated-vote	¶ 128,476


Table 6. Holding versus voting concentration for protocols with sufficient governance activity (12-month lookback). Voting HHI from Tally delegate voting power (on-chain) or Snapshot vote-weighted power (off-chain) or Solana on-chain governance (VSR lockup-weighted or delegated-vote per protocol design); N is unique-voter count (Snapshot rows; 12-month rolling window), top-1000 delegate sample (Tally rows), all-delegates count (Aave Tally all-delegates pull), or unique-voter count from Solana on-chain VoteRecord parsing (Solana rows). Ratio above 1.0 indicates delegation concentrates governance above token-holdings; below 1.0 indicates delegation distributes. Holding HHIs are post-exclusion per Table 3. Nine of thirteen sampled protocols amplify; four structural exceptions show dispersion (ENS 0.45x and GMX 0.87x per Sections 4.5.1 + 4.5.1.1; HNT 0.35x-0.53x sensitivity-range per Section 4.5.4 VSR-lockup-weighted design; JUP 0.057x per Section 4.5.4 delegated-vote + broad-airdrop-user-base; JUP is most-extreme dispersion outlier in the cross-section). Curve veCRV + Balancer veBAL (15x and 21x respectively) form a distinct ve-token class discussed separately in Section 4.5.2.

† Compound and Arbitrum voting HHI sourced from Snapshot (n = 138 and n = 5,466 unique voters respectively in the 12-month rolling-window measurement period) where Snapshot’s active-voter pool exceeded Tally’s top-100-delegate sampling. AAVE, Uniswap, Optimism, GMX, and ENS use Tally (EVM-side delegate registry); DIMO, WeatherXM, and Lido use Snapshot (Tally is not deployed for these EVM protocols). JUP, HNT, and DRIFT use Solana on-chain governance program parsing (JUP via Dune `jupiter_solana.govern_call_setvote` decoded namespace with per-signer max-weight aggregation; HNT and DRIFT via Helius RPC parsing of the respective SPL Governance program `VoteRecordV2` accounts) because Tally is not deployed for Solana protocols and on-chain governance state is the apples-to-apples comparable surface to Tally’s EVM delegate registries.

‡ GMX and ENS values from paginated Tally top-100 delegate pulls (May 2026). Snapshot date drift between the March 2026 baseline (used for the eight protocols with March 2026 voting data) and May 2026 (used for GMX and ENS) means rank-ordering across protocols is the durable inferential property, with absolute magnitudes subject to delegate-pool changes over multi-month windows.

§ Aave voting HHI 0.058 from Tally all-delegates pull (N = 150,911); the all-delegates depth establishes ARB at maximum 437,453-delegate depth as the deepest-N anchor for the methodology-floor sensitivity analysis. Four deeper-N sweeps (AAVE + COMP + UNI + ARB) produce sub-0.005 HHI shifts vs top-1000 baseline, confirming the top-1000 methodology floor.

¶ Solana voting HHI values from on-chain governance program parsing via Helius RPC and Dune decoded namespace: JUP via `jupiter_solana.govern_call_setvote` (per-signer $\max(\text{weight})$ aggregation; delegated-vote design; each token equals one vote); HNT and DRIFT via Voter Stake Registry (VSR) lockup-weighted vote design where vote weight is locked-amount multiplied by lockup-duration multiplier (HNT VSR saturates at four-year lockup; DRIFT VSR uses analogous lockup multiplier; both classified as VSR-lockup-weighted per DRIFT cycle source memo). HNT presented as sensitivity range under VSR position-state reconstruction with 14.8 percent closed-position rate bracketing the bounds. The within-Solana sample (N=3) is discussed in Section 4.5.4 within-Solana governance heterogeneity subsection; JUP at 0.057x ratio is the most-extreme dispersion outlier in the cross-section.



Source: 2026-05-22 multi-source PCA-symmetric extension cascade per EC-2026-05-22-B2-F-B2-15-DRIFT-VSR-Classification.
Snapshot rows: all proposals in 12-month rolling window 2025-05-22 to 2026-05-22. Tally rows: top-1000 (§: Aave all-delegates N=150,911).
Solana rows: JUP delegated-vote per `jupiter_solana.govern_call_setvote`; HNT and DRIFT VSR lockup-weighted (HNT shown as 0.026-0.039 midpoint).
N=13 sample: 9 amplifiers (range 1.6x DRIFT to 9.9x Compound); 4 dispersers (ENS 0.45x; GMX 0.87x; HNT 0.35-0.53x; JUP 0.057x most-extreme).

Figure 4. Raw holding HHI (lower bar) vs. delegation-adjusted voting HHI (upper bar) for ten Tally- and Snapshot-sourced protocols, sorted by amplification ratio. Reading the figure: the ratio label at the right of each pair gives voting HHI divided by holding HHI; ratios above 1.0 indicate delegation amplifies concentration (the upper bar is longer than the lower bar; green

annotation), ratios below 1.0 indicate delegation distributes it (red annotation). Eight of ten EVM/EVM-like protocols sampled show amplification, with ratios ranging from 2.7x (Uniswap) to 9.9x (Compound); two protocols (ENS at 0.45x and GMX at 0.87x) show dispersion (Section 4.5.1 + Section 4.5.1.1). The N=13 cross-section (with three Solana-native protocols: JUP delegated-vote 0.057x dispersion; HNT VSR 0.35x-0.53x dispersion; DRIFT VSR 1.6x amplification per Section 4.5.4 within-Solana heterogeneity discussion) shows nine of thirteen protocols amplify with four structural exceptions (ENS + GMX + HNT + JUP; JUP at 0.057x is the most-extreme dispersion outlier). Figure values: Snapshot rows aggregate unique voters across all proposals in the 12-month rolling window 2025-05-22 to 2026-05-22; Tally rows reflect top-1000 delegate samples per protocol.

Seven DeFi protocols sampled show varying delegation outcomes: JUP (0.057x, most-extreme dispersion outlier per Section 4.5.4 Solana delegated-vote with broad-airdrop-user-base mechanism), GMX (0.87x, disperses; methodology-sensitive crossover from amplification 1.2x at top-100 sample to dispersion at top-1000 sample; Section 4.5.1.1), DRIFT (1.6x, mild amplifier and the lowest amplifier in the cross-section; Solana VSR-lockup-weighted per Section 4.5.4), Uniswap (2.7x), Aave (4.4x), Lido (6.3x), and Compound (9.9x; the highest in the cross-protocol sample). The Uniswap result is worth examining in detail. Earlier analyses, using a holding HHI of 0.032, produced an amplification ratio of 0.84x and were cited as evidence that Uniswap's mature delegate ecosystem can disperse governance power. That earlier HHI included the canonical UNI burn address (0x000...000dead), which holds 102.5 million UNI (11.3 percent of supply) inert and cannot participate in governance under any condition: it has no delegate, no votes, and no path to ever cast one. After excluding this burn address, the post-exclusion holding HHI is 0.010, and the amplification ratio is 2.8x. Uniswap's delegate program concentrates voting power above its post-exclusion holding baseline, in line with every other protocol sampled. The pre-exclusion result was driven by including a wallet that structurally cannot vote in the holding denominator.

Both infrastructure protocols sampled show delegation amplification when holding HHIs are computed with the same exclusions applied to Table 3. Optimism's voting HHI (0.033) is 3.6 times its post-exclusion holding HHI (0.009). Arbitrum's voting HHI (0.038) is 3.1 times its holding HHI (0.012). Both protocols have invested in delegate-program infrastructure (published delegate platforms, recurring delegate compensation, active recruitment of long-form policy positions), and both produce amplification ratios in the 3.1x to 3.6x range, which sits below the cross-protocol mean. Delegate-program investment does not differentiate the two protocols on amplification magnitude. The institutional layers built around the delegate programs differ substantially and may differentiate outcomes that HHI cannot capture, as discussed in the next paragraph.

Optimism's Token House shows 3.6x delegation amplification, below the cross-protocol mean but consistent with the universal direction of effect. The Citizens' House moderates Optimism's governance in practice, though it does not enter the token-weighted HHI calculation. The Citizens' House is a separate governance body weighted by reputation rather than tokens, with formal veto power over economic parameters proposed by the Token House and a veto threshold that lowers when multiple stakeholder groups align (Optimism, 2024). Because reputation weights are not token weights, the Citizens' House provides accountability that the amplification ratio alone cannot capture.

Lido's 6.3x amplification catalyzed a parallel institutional response. The "Dual Governance" redesign grants stETH holders formal veto power over LDO governance proposals, with a programmatic "rage quit" exit mechanism as the ultimate accountability lever. The reform precedes the holding-HHI recomputation reported here and was justified at the time by the amplification ratio measured from a curated active-governance holder subset ($N = 189$; minimum balance threshold approximately 414K LDO). Under the universal top-1000 holder methodology applied in this manuscript with the full protocol-controlled-address audit (Section 3.9; Bybit and Binance custody addresses excluded as Class 5 CEX custody), Lido's amplification ratio is the most extreme in the sample, and the direction (concentration above stake) is consistent with the predominant-amplification pattern documented across nine of thirteen protocols.

Both responses share the same shape. Token-weighted delegation concentrates voting power above stake holdings across the sample, and the institutional reforms operate by adding an accountability layer outside the token-weighted system rather than redesigning the delegation mechanism itself. Both protocols with named accountability reforms (Optimism, Lido) are protocols where token-weighted delegation shows positive amplification, consistent with reading these reforms as responses to the predominant-amplification pattern that this section documents.

The delegation-adjusted analysis reported in this section extends recent empirical work by Fritsch, Müller, and Wattenhofer (2024), who apply Nakamoto coefficients and Gini measures to voting power in Compound, Uniswap, and ENS governance systems and document extreme concentration in delegate-level voting power. The cross-protocol comparison reported here spans thirteen protocols across DeFi (Compound, Aave, Uniswap, GMX, Lido), DePIN (DIMO, WeatherXM), and infrastructure (Optimism, Arbitrum, ENS). Amplification factors range from 1.6x (DRIFT) to 9.9x (Compound) across the nine amplifying protocols; ENS at 0.45x; GMX at 0.87x; HNT at 0.35x-0.53x; and JUP at 0.057x are the four structural exceptions (ENS: mature delegate program where holders systematically delegate to a broad community-delegate set; GMX: methodology-sensitive crossover per Section 4.5.1.1; HNT: Solana VSR lockup-weighted design per Section 4.5.4; JUP: Solana delegated-vote with broad airdrop user base per Section 4.5.4; JUP at 0.057x is the most-extreme dispersion outlier in the cross-section). Within the sample, the 1.6x to 9.9x amplification range is distributed across sectors with no clear concentration in any one; sector membership does not appear to be associated with how strongly delegation concentrates power.

The pattern is consistent with the broader claim of this paper: design choices within each delegate program (delegate selection criteria, compensation structure, voting threshold rules) determine the magnitude of governance concentration through delegation, while the direction of the effect (concentration above holdings) is the same across sector and governance venue. Whether a protocol concentrates governance power through delegation does not depend on whether it is DeFi, DePIN, or infrastructure. How much it concentrates is shaped by design choices made at the level of the delegate program itself. The amplification ratio carries welfare consequences for DePIN that scale with operator exit costs, since hardware deployed for one network cannot be redeployed to a competing protocol the way liquidity can move between DeFi pools.

4.5.1 ENS as institutional-design counterexample

ENS is one of two protocols in the cross-section where delegation empirically disperses voting power below the holding-HHI baseline (ratio 0.45x; voting HHI 0.022 vs holding HHI 0.049, the latter computed after the ENS DAO Wallet and Treasury PCA exclusions codified in Section 3.8). The dispersion case is informative for the framework's institutional-design thesis because it provides an existence-proof that delegation outcomes are design-tractable rather than mechanism-inevitable. The ENS delegate program funds approximately 100 active delegates (including the ENS DAO ecosystem grants for delegate work), and the top-100 delegate voting-power distribution is structurally broader than the holding distribution: top-1 post-exclusion holding share is 16.8% while top-1 voting-power share is 12.1% (fireeyesdao.eth, a community-elected delegate). Holders systematically delegate to a broader community-delegate set rather than self-delegating or delegating to whales.

The ENS + GMX + HNT + JUP dispersion findings contrast with the nine protocols where delegation amplifies: in those protocols, the top-delegate identity is typically either the protocol foundation itself, a staking-aggregation contract that aggregates underlying staker voting power, or an institutional delegate firm (see Section 4.5.3). The structural pattern across the thirteen-protocol sample is consistent with the framework's thesis that institutional design at the delegate-program level governs the direction of delegation's effect on voting concentration. The generalizable design recommendation: delegate-program architectures that incentivize broad community-delegate distribution (recurring delegate compensation; public delegate platforms; active recruitment; operational support for non-whale delegates) can produce delegation-mediated dispersion of voting power, contrary to the default-amplification pattern documented across most token-governance protocols.

4.5.1.1 GMX as second structural-exception counterexample

GMX is the second protocol in the cross-section where delegation empirically disperses voting power below the holding-HHI baseline (ratio 0.87x; voting HHI 0.057 vs holding HHI 0.065, the latter computed after the universal burn-rule exclusion plus GMX-specific Binance-custody exclusions codified in Section 3.8). The dispersion case is methodology-sensitive: at a top-100 delegate sample GMX appears to amplify at 1.2x, while at the canonical top-1000 sample the broader delegate distribution surfaces voting-power-share tails less concentrated than the holding-share tail at the same sampling depth. The crossover from amplification (1.2x at $N = 100$) to dispersion (0.87x at $N = 1000$) indicates that GMX top-100 amplification readings are sampling-depth artifacts rather than stable properties of the underlying delegate-program design; the top-1000 sample better captures the actual voting-power distribution.

The finding strengthens the framework's institutional-design thesis: GMX's delegate program is observably broader than its holder concentration when measured at appropriate sampling depth, consistent with the mature-delegate-program pattern documented at ENS (Section 4.5.1) at a different concentration level. The two structural-exception cases (ENS at 0.45x; GMX at 0.87x) together preserve the dispersion pattern as informative counterexample to the predominant-amplification thesis rather than refuting the broader argument that token-weighted delegation predominantly concentrates voting power above holdings.

4.5.2 Secondary extension: vote-escrowed governance

Two protocols using vote-escrowed governance (Curve veCRV; Balancer veBAL) form a distinct extreme-amplification class that operates through structurally different mechanisms than the delegate-based amplification documented in the Table 6 sample. Curve's veCRV system locks CRV for up to four years; lock-duration multipliers convert locked CRV into proportional veCRV voting power, with longer locks weighted more heavily. Convex (the dominant meta-governance protocol for Curve) aggregates user-deposited CRV into a single locked position that votes as a meta-governance layer; Convex holds approximately 48 percent of the top-50 historical veCRV deposit-weighted share per the direct cumulative-deposit computation (Section 4.5.2; supplementary file `veCRV_voting_concentration`). The resulting veCRV voting concentration (HHI 0.26 from cumulative-deposit aggregation across top-50 lockers) is approximately 15 times the raw post-Convex- exclusion CRV holding HHI of 0.014.

Balancer's veBAL system exhibits a structurally identical pattern: Snapshot voting on Balancer's governance space (12-month lookback; 36 active voters) yields voting HHI 0.626 against holding HHI 0.030, a 21x amplification ratio. The top-1 active voter holds 78 percent of total voting power; the top-10 voters cumulatively hold 99.85 percent. The Aura Finance protocol (Balancer's Convex-equivalent meta-governance aggregator) plays the same role for veBAL that Convex plays for veCRV.

The veCRV and veBAL cases are not included in the Table 6 row set because the underlying methodology differs: Table 6 measures HHI of delegate voting power within token-governance delegate programs; the ve-token cases measure HHI of lock-duration-weighted CRV/BAL deposits, with meta-governance aggregation. The two mechanism classes produce concentration through different paths but converge on the same direction (amplification) with much larger magnitudes (15-21x versus the 2.7x-9.9x range observed in delegate-program-based amplification across the nine amplifying protocols). For practitioners designing token governance, the ve-token class is a separate design choice with substantially different concentration implications than standard delegate-program governance.

4.5.3 Foundation and aggregation-contract top-delegate pattern

Three of the thirteen Table 6 protocols (plus the two ve-token class protocols in Section 4.5.2) have their largest single voting block held by a protocol-controlled aggregation entity rather than an independent community delegate: Compound Foundation (21.5% of total delegated voting power in May 2026 Tally pull); Aave stkAAVE staking contract (21.9%; aggregates underlying staker voting power as a single delegate); Curve Convex (48% of top-50 deposit-weighted veCRV share per Section 4.5.2). The three Solana-native protocols added at Table 6 N=13 expansion (JUP, HNT, DRIFT; Section 4.5.4) do not exhibit the foundation-or-aggregation-entity top-delegate pattern at the Solana SPL Governance member-set layer; Section 4.5.4.1 documents the Drift Foundation proposal-authorship-versus-vote-weight decoupling pattern that further differentiates Solana governance topology from the EVM precedents described here.

The aggregation-entity pattern raises practical measurement concerns relative to independent-delegate-dominated protocols (Uniswap top-1 is an investor delegate; Optimism top-1 is a professional delegate firm; ENS top-1 is a community delegate). Aggregation entities may or may not represent the distributional interests of the underlying token-holders or stakers who

delegated to them; the principal-agent gap (Holmstrom, 1979; Tirole, 2001) is amplified by the aggregation layer. The PCA-symmetric exclusion analysis (Section 4.6) shows that predominant amplification holds even when aggregation entities are excluded; the empirical finding is robust to this methodological consideration, but the structural pattern itself remains a substantive feature of token-governance practice.

4.5.4 Within-Solana governance concentration heterogeneity

Three Solana-native protocols in the cross-section have on-chain voting HHI computable via proper-weight methodology: Jupiter (JUP) at 0.00546 via per-signer `max(weight)` aggregation on `jupiter_solana.govern_call_setvote` (delegated-vote design); Helium (HNT) at 0.0261 to 0.0394 (sensitivity range under Voter Stake Registry lockup-weight reconstruction; 14.8 percent closed-position rate brackets the bounds); and Drift Protocol (DRIFT) at 0.0833 via direct on-chain parsing of `VoteRecordV2` accounts in the DRIFT governance program (VSR lockup-weighted design analogous to HNT). The within-Solana spread (0.005 to 0.083; 15x range) sits within the broader within-EVM voting HHI spread (full range 0.022 to 0.556 across the EVM voting-HHI cross-section; range 0.022 to 0.089 if the WeatherXM thin-pool outlier with N=73 voters is excluded). The apples-to-apples within-Solana versus within-EVM comparison requires acknowledging that small-voter-pool protocols on both chains (WeatherXM N=73 on EVM; DIMO N=10 on EVM-Polygon; analogous small-pool Solana protocols outside the voting-HHI subset due to data availability) skew their respective ranges upward; the within-Solana sample does not include a comparable thin-pool outlier in the voting-HHI subset. Both Helium and Drift sit within the EVM-typical voting HHI range (HNT in the Arbitrum neighborhood at 0.0376; DRIFT in the Uniswap-to-Compound neighborhood at 0.0727 to 0.0889). Jupiter is the only Solana protocol exhibiting voting concentration order-of-magnitude below the EVM cross-section median.

The within-Solana heterogeneity is structurally organized by voting-weight mechanism design as the primary axis, with user-base composition and rewards-calibration as secondary axes. The primary driver: two of the three Solana protocols (Helium HNT; Drift DRIFT) use Voter Stake Registry (VSR) lockup-weighted vote design, where vote weight is locked-amount multiplied by a lockup-duration multiplier (HNT VSR saturates at four-year lockup; Drift VSR uses an analogous lockup multiplier within its governance program). Both VSR-using Solana protocols cluster at EVM-typical voting HHI (HNT 0.0261 to 0.0394; DRIFT 0.0833). Jupiter JUP, by contrast, uses delegated-vote weight (each token equals one vote at vote time; no lockup multiplier), and Jupiter is the only Solana protocol exhibiting voting concentration order-of-magnitude below the EVM median. The VSR design choice penalizes short-lockup voters and rewards long-lockup whales (whales can afford to lock larger positions for longer durations), raising concentration relative to per-token weighting; delegated-vote design flattens the vote-weight distribution within the active voter pool. Secondary drivers within the design classes: user-base composition (Jupiter's 128,476-voter pool was built by a broad launch airdrop, producing a retail tail that flattens distribution within the delegated-vote class); rewards calibration (Jupiter's Active Staker Rewards directly compensate voting participation at a high engagement rate; Helium's VSR rewards lockup-and-vote but the lockup-weight calculus advantages whales; Drift lacks a comparable direct vote-rewards mechanism).

The pattern matters for chain-comparative governance analysis. A categorical claim that Solana governance is less concentrated than EVM governance is not supported by the proper-weight

Solana sample; the apparent low-concentration finding from vote-count-proxy methodology on Helium was an artifact superseded by lockup-weight reconstruction. Chain context alone is insufficient to produce low-concentration governance, as HNT and DRIFT demonstrate by clustering in EVM-typical concentration ranges. JUP's design choices (airdrop engagement; Active Staker Rewards; retail orientation) are protocol-design-level decisions that produce the outlier; these decisions are enabled by Solana's low-cost voting environment, but whether they would replicate at scale on a higher-fee chain is untested because JUP is Solana-native. Protocol design and chain context co-vary for chain-tied protocols, which limits clean protocol-level versus chain-level decomposition. Section 4.5.4.1 documents the Drift Foundation proposal-authorship versus vote-weight decoupling pattern that further differentiates Solana governance topology from EVM precedents.

4.5.4.1 Drift Foundation proposal-authorship and vote-weight decoupling

Drift Protocol exhibits a governance pattern with no close analog in the EVM-protocol cross-section. A single Drift Foundation submitter address (4CL25...wPiIF on Solana) proposed 11 of the 13 DRIFT Improvement Proposals (DIPs) over the 12-month window, performing the agenda-setting function for protocol governance. The same address holds only 0.18 percent of total DRIFT vote weight in the voter-pool over the same window. Proposal-authorship and vote-weight are structurally decoupled: the Foundation submits proposals, the token-weighted voter pool decides, and the two functions are not co-located in the same address.

This decoupling pattern is methodologically informative. EVM-protocol top delegates are typically both agenda-setters and vote-weight-holders (Compound Foundation as the Tally top delegate in Section 4.5.3 both proposes and votes from a foundation-controlled wallet at 21.5 percent share). The Drift Foundation structure separates these functions explicitly, with the Foundation address operating in a low-vote-weight agenda-setter role while the broader token-holder community carries the decision-rights. The DRIFT Squads V4 multisig PCA (9Wiivy8...) does not appear in the Drift governance voter set, confirming the Solana governance topology bifurcation documented in Section 4.6 (on-chain SPL voting member set is structurally separate from off-Realms custody multisigs).

4.5.5 Secondary extension: cross-protocol delegate overlap

Per-protocol holding HHI and per-protocol voting HHI both measure concentration within a single protocol's governance surface. A distinct concentration dimension, made visible only by cross-protocol aggregation, is the institutional-governance-investor pattern. The Snapshot voter-pool extension to 13 protocols aggregated 11,077 unique voter addresses across the cross-section (UNI, COMP, LDO, DIMO, WXM, ARB, BAL, GTC, MPL_SYRUP, POKT, AAVE, ENS, GMX). Of these, 1,801 voters (16.3 percent) participate in 2 or more protocols' Snapshot governance. A small institutional-governance-investor class (approximately 10 to 50 addresses) participates in 6 or more protocols' delegate programs simultaneously.

The top three institutional cross-protocol delegators collectively exert 17.0 percent combined Snapshot voting power across the L1 DeFi suite. PGov (0x3fb19771...; pgov.eth) participates in 7 DAOs (ARB, COMP, ENS, GMX, GTC, LDO, UNI) at 10.16 percent combined summed share. Tane Governance (0xb79294d0...; tanegov.eth) participates in 6 DAOs at 4.83 percent. Arana Digital (0x0579a616...; aranadigital.eth) participates in 6 DAOs at 2.04 percent. These

are professional governance firms whose business model is providing cross-protocol delegation services to token-holders who do not vote directly; the firms aggregate delegated voting power across multiple protocols simultaneously.

The pattern is invisible to per-protocol concentration metrics. The predominant-amplification finding (Section 4.5) addresses delegation amplification within a single protocol (delegated voting power exceeds post-exclusion holding HHI); it does not address whether the same set of institutional delegates appears across multiple protocols. Cross-protocol institutional concentration is a candidate fifth concentration axis alongside the four already in the paper's scope: holder concentration (Section 4.2), voter concentration (Section 4.5 Table 6), delegate-class concentration (Section 4.5.3 foundation-and-aggregation-contract pattern), and burner concentration (Section 4.6 universal burn-rule application). Token-holders who delegate to professional firms acquire delegated representation across multiple protocols simultaneously; the firms' decisions affect governance outcomes in all protocols where they hold delegated weight, creating correlated voting patterns invisible to per-protocol Mann-Whitney or Cohen's d analysis. The Tally-side cross-protocol overlap network is deferred to future work.

4.6 Robustness: Leave-One-Out Sensitivity

The cross-sectional sample size ($N = 30$ for the DePIN-vs-DeFi sector contrast; $N = 22$ for the subsidy excluding Livepeer test) bounds the strength of inference. Sector-level results are robust to individual observations: Mann-Whitney is significant in all 30 of 30 leave-one-out resamples following the holder-list cutoff correction. Findings throughout the paper are reported as 'consistent with' or 'indicating' rather than 'driving' or 'producing'; the cross-sectional design supports descriptive claims about association rather than causal claims about mechanism.

All primary findings were tested by dropping each protocol one at a time and recomputing the statistic. The insider-retention correlation is fully robust. The DePIN-DeFi sector difference is fully robust to single-protocol exclusion (significant across all 30 of 30 leave-one-out iterations). The subsidy correlation is entirely driven by Livepeer. Details follow.

The insider concentration relationship (Spearman $\rho = 0.48$, $p = 0.003$, $N = 37$) is robust: it remains significant across all 37 single-protocol exclusions, with ρ ranging from 0.45 to 0.55. Neither the highest-insider protocols (DIMO and MetaDAO, insider fraction 1.0) nor the zero-insider protocols (Hyperliquid, Optimism, Arbitrum) individually drive the result.

Secondary size and retention checks sharpen the interpretation. Larger protocols by fully diluted valuation have fewer insider wallets in their current top-10 holder set (Pearson $r = -0.36$, $p = 0.032$, $N = 36$; Spearman $\rho = -0.31$, $p = 0.063$), and the market-capitalization specification has the same direction at borderline significance (Pearson $r = -0.34$, $p = 0.050$, $N = 33$). Balance-retention specifications are not significant, indicating a composition-shift interpretation rather than a retained-balance result. Net-deflationary protocols also show no statistically significant difference in insider retention by count or balance (Mann-Whitney $p = 0.57$ and $p = 0.78$, respectively). These checks are secondary, but they support the paper's main distinction between launch allocation and current holder composition.

The DePIN-versus-DeFi sector difference (Mann-Whitney $p = 0.020$, Cohen's $d = 0.94$) is robust to sample composition. The test remains significant in all 30 of 30 leave-one-out iterations (per-iteration p range: 0.008 to 0.034; Cohen's d range: 0.84 to 1.10; per-iteration LOO data (see

Supplementary Figure S5b in Supplementary File S5 for the LOO forest plot) covers the per-iteration forest plot with bootstrap 95% CI on the headline d). Seven protocols, when individually dropped, push the p-value to 0.034: Compound, Uniswap, Lido, Hyperliquid, WeatherXM, Livepeer, and io.net. No single protocol uniquely drives the result. A two-sided permutation test (100,000 random reassignments of sector labels) yields $p = 0.012$, consistent with the direction not being an artifact of the asymptotic approximation. A non-parametric bootstrap on the sector mean difference (10,000 resamples) yields a 95% percentile interval of $[+0.012, +0.069]$ HHI points, strictly excluding zero; the Cohen's d 95% percentile interval is $[0.32, 1.68]$, strictly above the small-effect threshold and including the conventional large-effect threshold ($d = 0.8$) within the upper portion of the interval. We characterize this finding as robust under sample composition, resampling, and rank-based and parametric tests.





Figure 5. Leave-one-out forest plot for the DePIN-versus-DeFi sector contrast. Each row reports Cohen's d (point estimate) for the Mann-Whitney sector test recomputed after dropping the named protocol from the 30-protocol sector sample ($N = 15$ DePIN + 15 DeFi). All 30 iterations remain statistically significant at $p < 0.05$ (per-iteration p range 0.008 to 0.034); the Cohen's d point estimates span 0.84 to 1.10 across iterations. The bootstrap 95 percent percentile interval on the headline d (10,000 resamples; full-sample estimate) is $[0.32, 1.68]$, strictly above the small-effect threshold and including the conventional large-effect threshold ($d = 0.8$). No single protocol drives the sector contrast; the finding is structurally robust to single-observation perturbation.

The subsidy-to-concentration correlation (Pearson $r = 0.59$, $p = 0.003$, $N = 23$) is driven primarily by a single protocol. Livepeer's subsidy ratio of 88.5x is a 3.5-sigma outlier on the subsidy axis, and its HHI (0.20) is the second-highest in the subsidy subsample. Excluding Livepeer alone reduces the correlation to a clean null at $r = 0.07$ ($p = 0.77$). The remaining 22 protocols show no detectable linear association, indicating that the headline subsidy result is entirely Livepeer-driven. The Spearman rank correlation on the full subsidy sample is also insignificant ($\rho = 0.20$, $p = 0.35$), and a log transformation eliminates the linear relationship (r

= 0.27, $p = 0.22$, $N = 23$). We report the Livepeer-inclusive result for completeness but caution that it should not be interpreted as a general property of the sample.

A robustness check using the Token Terminal subsidy_ratio ($N = 20$ protocols with non-null Token Terminal revenue and incentives data) is consistent with the null at cross-sector Pearson $r = 0.12$ ($p \approx 0.61$); a broader merged sample combining Token Terminal data with on-chain emission/fee ratios ($N = 26$ across both methodologies; $N = 23$ protocols with non-zero subsidy under either metric) yields $r = 0.59$ inclusive of Livepeer and $r = 0.07$ excluding Livepeer. The cross-sector methodology aggregates by preference: Token Terminal subsidy_ratio is used where available ($N = 20$ protocols with non-empty TT data); on-chain emission/fee subsidy_ratio_onchain is used as fallback where TT data is unavailable but on-chain data is present (additional $N = 6$ protocols); the union yields the merged $N = 26$ sample. The “non-zero either metric” sub-sample ($N = 23$) filters to protocols where either $TT > 0$ OR $OC > 0$; this excludes 3 protocols with TT subsidy_ratio = 0 explicitly (no subsidy under either methodology). The methodology is documented in this footnote pre-emptively to facilitate reviewer reproduction of the cross- sector subsidy correlation values, qualitatively replicating the main result that subsidy and HHI co-vary only through the Livepeer outlier.

A multivariate specification with both subsidy ratio and DePIN sector indicator on the 23-protocol non-zero-subsidy sample (HC3 robust standard errors, a variance estimator that does not assume equal error variance across observations) further reinforces the Livepeer-driven interpretation: with all 23 protocols, the DePIN sector indicator is significant ($p = 0.016$) while subsidy is not ($p = 0.36$; Adj R-squared = 0.47); excluding Livepeer, the subsidy coefficient is strongly non-significant ($p = 0.94$) while the DePIN sector indicator is significant ($p = 0.007$), indicating that the apparent subsidy-HHI relationship is fully absorbed by sector membership absent the single extreme observation. Per-sector results in [Supplementary File S8](#).



Source: Author calculations from Dune Analytics and Token Terminal (March 2026). N=25.

Figure 6. Subsidy ratio (emissions/revenue) vs governance concentration (HHI) for 23 protocols with on-chain subsidy data, with per-protocol labels. Livepeer (diamond, 88.5x) is a 3.5-sigma outlier that alone drives the Pearson correlation: $r = 0.59$ (solid line, full sample) drops to $r = 0.07$ excluding Livepeer (line omitted to communicate null result honestly; $p = 0.77$). The Token Terminal robustness sample ($N = 20$, [Supplementary File S8](#)) is consistent with the null at $r = 0.12$. The cross-sector association drops to a clean null when Livepeer is excluded, indicating the headline subsidy result is sensitive to a single high-leverage protocol.

Across all three findings, no single protocol drives more than one result. Livepeer appears both in the set of sector-test single-drop protocols that push Mann-Whitney p toward 0.05 and as the sole protocol that, when dropped, eliminates the subsidy correlation; however its sector influence is shared equally among seven protocols, whereas its subsidy influence is singular.

The predominant-amplification finding in Section 4.5 also requires a robustness check on the voting-HHI methodology. The Table 6 voting HHI follows established delegation-HHI literature (Fritsch et al., 2024) by including all delegates with non-zero voting power, regardless of whether the delegate is a protocol-controlled entity. Applying protocol-controlled-address exclusion symmetrically (excluding foundation, treasury, and aggregation-contract delegates from voting HHI as we exclude them from holding HHI) requires careful methodology because PCA addresses appear in three structurally distinct surface classes: (a) holder-side custody addresses (treasury safes, vesting multisigs, foundation cold wallets, staking-aggregation contracts, CEX hot wallets, burn destinations) captured in `exclusions_log.csv`; (b) signer-side governance-operation addresses (foundation operational signing wallets distinct from

custody addresses; for example Compound Foundation 0xb06df4dd . . . , Balancer DAO multisig 0xad9992f3 . . . , Kevin Owocki / Gitcoin Maintainer 0x00de4b13 . . . , WXM Hedgex-vested insider 0x96d8901a . . .) captured in the sibling canonical store `exclusions_log_signer.csv` introduced this revision; and (c) off-Realms custody multisigs for Solana protocols (Drift Foundation Squads V4 multisig 9Wiivy8 . . . ; Helium Foundation treasury AguTdj . . .) which do not appear in either the on-Realms voter set or in Snapshot voter pools because Solana PCAs structurally separate custody from governance signing.

Applying symmetric PCA exclusion across all three surface classes for the cross-section yields several substantive results.

For four Tally-sourced protocols at all-delegates depth (AAVE N=150,911; COMP N=18,627; UNI N=48,707; ARB N=437,453), PCA-symmetric voting HHI values shift modestly from the top-1000 sample baseline: AAVE 0.0357 to 0.0324 (delta minus 0.0033); COMP 0.0516 unchanged across the top-1000-to-all-delegates sweep (three additional PCAs detected at deeper N have negligible voting share); UNI 0.0267 unchanged; ARB 0.0340 to 0.0320 (delta minus 0.002). All four sweeps produce HHI shifts below 0.005. ARB at the maximum 437,453-delegate depth serves as the deepest-N anchor for reviewer-defensibility against the methodology-asked surface. The remaining six Tally-sourced protocols (OP, ENS, GMX, GTC, W, DIMO) retain top-1000 sample-depth values as the methodology-floor canonical, empirically validated by the four deeper-N sweeps. The top-1000 sample is robust because PCAs concentrate at the top of the distribution and tail delegates have share-squared contributions below ten to the negative four.

For Snapshot voter pools across 13 protocols (UNI, COMP, LDO, DIMO, WXM, ARB, BAL, GTC, MPL_SYRUP, POKT, plus cross-source AAVE, ENS, GMX), holder-side PCA-address-based symmetric exclusion detects zero PCA voters across 12,878 voter rows. Snapshot voter pools (off-chain message signers) are structurally disjoint from the holder-side PCA-address taxonomy. However, functional-PCA voters operate via signer-side addresses that are not captured in `exclusions_log.csv` but ARE substantial in voting weight: Compound Foundation (25.01 percent of COMP Snapshot voting weight); Balancer DAO multisig (51.95 percent of BAL); Kevin Owocki / Gitcoin Maintainer (27.93 percent of GTC); WXM Hedgex-vested insider (74.04 percent of WXM). The original DIMO Foundation candidate cluster (rank-1 plus rank-2; 58.83 percent combined) was registered as a candidate protocol in the initial expansion but subsequently REFUTED via Polygon Blockscout MCP funding-source verification (neither voter holds substantial DIMO tokens; the cluster represents active community delegates participating in a thin N=10 voter pool, not functional PCAs); the DIMO entries were removed from `exclusions_log_signer.csv`. The `exclusions_log_signer.csv` canonical store registers the four confirmed signer-side functional PCAs (COMP + BAL + GTC + WXM).

A methodologically substantive finding from the Snapshot signer-side PCA-symmetric recompute is that the exclusion direction is NOT monotonic in the concentration-reduction direction. For three of the four Snapshot-sourced protocols with confirmed signer-side functional PCAs, exclusion produces the expected concentration-reduction direction: COMP HHI shifts from 0.0889 to 0.0468 (delta minus 0.0421); BAL from 0.3045 to 0.1499 (delta minus 0.1546); WXM from 0.5561 to 0.1184 (delta minus 0.4377). For the remaining protocol, signer-side PCA exclusion INCREASES voting HHI: GTC shifts from 0.1792 to 0.1948 (delta plus 0.0156). For comparison, the initial DIMO Foundation candidate exclusion (before the candidate cluster was

refuted via Blockscout MCP verification) would have produced a 0.2280 to 0.3242 shift (delta plus 0.0962); this computation is reported for completeness. The reverse-direction pattern arises from a smaller-denominator effect: when the excluded PCA holds a dominant share of total voting weight (GTC top-1 at 28 percent; DIMO candidate cluster at 59 percent combined pre-refutation), the remaining voters become MORE concentrated relative to each other under the smaller post-exclusion total. The pattern is most extreme in small voter pools (DIMO N equals 10 voters); larger pools dilute the effect. The methodological implication is that PCA-symmetric exclusion is a substantive direction-of-effect test rather than a mechanical concentration-reduction adjustment: when exclusion increases HHI, the finding indicates that PCA concentration was masking a more concentrated underlying distribution among non-PCA voters, which is itself a substantive structural feature of the protocol's governance.

For Solana governance (three protocols at proper-weight methodology), the apples-to-apples comparable voting HHI values are JUP 0.00546 (Jupiter; via per-signer max(weight) on `jupiter_solana.govern_call_setvote`); HNT 0.0261 to 0.0394 (Helium; sensitivity range via VSR position-state reconstruction; 14.8 percent closed-position rate documents the bounds); DRIFT 0.0833 (Drift Protocol; via Solana RPC parsing of the DRIFT governance program `VoteRecordV2` accounts). The DRIFT Squads V4 PCA address does not appear in the DRIFT voter set, confirming the Solana structural separation between custody and governance signing documented in the Realms governance survey. JUP is the only Solana protocol exhibiting voting concentration order-of-magnitude below the EVM cross-section median; HNT and DRIFT cluster within EVM-typical voting HHI range. Section 4.5.4 documents the within-Solana governance heterogeneity; Section 4.5.4.1 documents the Drift Foundation proposal-authorship-versus-vote-weight decoupling pattern; Section 4.5.5 documents the cross-protocol institutional-governance-investor pattern.

For Realms-resident Solana DAOs (DRIFT, HNT parent, HNT MOBILE and IOT sub-DAOs), the Realms public surface exposes participation-count proxy not VSR-weighted voting power without a connected wallet; voting-side HHI in the Tally-equivalent specification cannot be computed from the Realms DOM surface alone. Direct on-chain SPL Governance program parsing (the approach used for DRIFT in this revision) is the methodological alternative. Five Solana protocols in the cross-section (HONEY, IO, META, RENDER, GRASS) host governance off-Realms and outside the Snapshot space, producing a documented gap in voting-side coverage for those protocols; holder-side concentration metrics from the Helius DAS API top-1000 holders pull remain the canonical Solana coverage for the gap-protocol subset.

The methodology-floor canonical for the voting-HHI cross-section is top-1000 sample depth for Tally + Snapshot, plus proper-weight methodology for VSR-using protocols (Helium VSR; Jupiter governance; Drift governance); the predominant-amplification finding (Section 4.5) is robust to: (a) top-1000 versus all-delegates sampling depth on Tally (empirically validated by 4-of-10 protocols at all-delegates depth with sub-0.005 HHI shifts); (b) PCA-symmetric exclusion via holder-side plus signer-side canonical stores on Snapshot (under both monotonic and non-monotonic direction-of-effect outcomes); (c) proper-weight versus vote-count methodology on VSR-using protocols (the vote-count proxy understates concentration by 5x to 7.5x for HNT per the multi-source extension EC log). Full per-protocol detail and signer-side PCA case interpretations in [Supplementary File S12](#).

The sector-contrast robustness extends to the PCA-classification specification choice. Table 7 reports the Mann-Whitney sector test recomputed under five progressively stricter alternative exclusion specifications (Specs A through E; full per-specification narrative in Supplementary File S10). The directional claim (DePIN HHI > DeFi HHI in central tendency) holds under every specification (Cohen’s *d* positive in all five). Statistical significance at the conventional 0.05 threshold requires the canonical 5-class typology (Spec A) or near-equivalents (Spec B drops Class 5 CEX custody and retains significance); progressively reducing exclusion strictness absorbs non-vote-eligible custody mass into the holding distribution and exhausts the sector signal at Specs D and E. The PCA exclusion methodology applied throughout the paper is therefore load-bearing for the inferential claim of statistical significance; the directional claim is robust to specification choice.

Spec	Description	Excluded *	* <i>p</i> **	* <i>d</i> **
A	Canonical 5-class (all 5 classes excluded)	125 / 36 prot	0.029	0.94
B	Drop Class 5 (CEX custody retained)	106	0.039	0.84
C	Drop Class 4 + 5 (bridge + CEX retained)	98	0.115	0.60
D	Drop Class 3 + 4 + 5 (only Classes 1 + 2 excluded)	86	0.610	0.17
E	Drop Class 2 + 3 + 4 + 5 (only burn destinations excluded)	10	0.512	0.35

Table 7. PCA-classification robustness: Mann-Whitney sector test under five alternative exclusion specifications applied to the $N = 30$ DePIN-vs-DeFi sector contrast. Specs A and B retain conventional statistical significance ($p < 0.05$) with large Cohen’s *d*; Spec C produces a medium-effect contrast that fails the 0.05 threshold at $N = 30$; Specs D and E (minimal exclusions) absorb so much non-vote-eligible signal into the holding distribution that the sector-contrast statistical power is exhausted. Cohen’s *d* is positive across all five specifications: the directional cross-protocol pattern is not an artifact of specific PCA-exclusion choices. Per-specification narrative interpretation and alternative-classification rationale in Supplementary File S10.

Other extended robustness analyses are documented in supplementary materials rather than reproduced in full here: Theil and Atkinson concentration metrics (S9), Shapley-Shubik and Banzhaf power indices (S11), PCA-symmetric voting HHI and signer-side functional PCA cases (S12), multivariate subsidy models with sector controls (S8), and the canonical statistics ledger (S0). Profitability, size, and insider-retention specifications are reported in the §4.6 secondary-checks paragraph above. Rank ordering across concentration metrics is stable (HHI-Theil Pearson $r = 0.77$ on the post-exclusion sample).

Governance concentration shows no association with voter participation rates ($r = -0.10$, $N = 13$ across 13 protocols with Snapshot or Governor governance data; see Supplementary Figure S5a in Supplementary File S5 for the HHI-vs-participation scatter), consistent with the Olson (1965)

collective action prediction that participation decisions reflect rational abstention and diffusion of responsibility rather than ownership structure.

5 Discussion

5.1 Measurement Implications

The results change how governance concentration should be audited. Launch allocation tables are not enough: initial insider allocation is not associated with post-distribution HHI in this cross-section. Audits should instead begin with current holder composition, explicit PCA exclusions, insider wallet retention among top holders, and delegated voting power.

Three practical implications follow. First, holder-list HHI should be computed only after excluding addresses that cannot vote as independent governance actors. Second, delegation analysis should compare voting-HHI with post-exclusion holding HHI, because delegated power usually differs from token holdings. Third, concentration audits should treat vote-escrow systems and delegation systems as different mechanism classes rather than aggregating them under one generic governance label.

5.2 Governance Design Implications

The empirical findings do not prove that any one design is sufficient for deconcentration, but they identify design levers that merit closer evaluation. Two counterexamples operating in different mechanism domains illustrate that concentration variation is design-tractable rather than mechanically inevitable: Anyone Protocol shows that low DePIN holding concentration is possible despite sector-level concentration risk, and ENS shows that delegation can disperse voting power below the holding baseline under a mature delegate program. GMX, HNT, and JUP extend the delegation-dispersion case under different mechanism designs, while Curve and Balancer show that vote-escrowed governance can amplify voting concentration far above raw token holdings.

For practitioners, the immediate design lesson is to measure current governance power rather than infer it from launch documents. Protocols should publish PCA-adjusted concentration metrics, delegate-power concentration, and any vote-escrow amplification before asking users or regulators to accept decentralization claims.

5.3 Risks and Counterpoints

Several empirical risks deserve candid acknowledgment. First, concentration may be efficient in some protocols if high-stake holders are also the participants most informed about protocol maintenance. This paper does not resolve that normative dispute; it shows that concentration should be measured after PCA correction and voting-power adjustment before decentralization claims are made. Second, launch allocation is empirically uninformative about current concentration, so audits that rely on allocation documents alone will miss post-distribution accumulation. Third, voting-HHI depends on governance-platform data that can drift over time, making rank ordering more durable than exact point estimates.

Measurement remains incomplete. Informal influence, off-chain coordination, and legal or equity-holder obligations can affect governance without appearing in token-holder lists. The paper therefore treats HHI and voting-HHI as necessary but not sufficient measures of governance power.

5.4 Position in Literature

This paper contributes to the empirical DAO-governance literature by refining how concentration is measured. Prior work documents token inequality, voter concentration, participation gaps, and the decentralization illusion. B2 extends that literature by showing that concentration estimates change materially once protocol-controlled addresses are excluded, and by comparing holding HHI with voting-HHI across delegation systems.

The paper also contributes to DePIN tokenomics by separating governance concentration from subsidy dependence and launch allocation. DePIN protocols exhibit higher concentration than DeFi in the corrected cross-section, but the sector pattern coexists with low-concentration counterexamples such as Anyone Protocol and Hivemapper. This makes the paper a measurement contribution rather than a deterministic claim about DePIN architecture.

5.5 Summary of Findings

Across 40 protocols, naive holder lists are a weak basis for governance claims. PCA exclusion corrects large measurement artifacts, with median HHI inflation of 2.3x and a maximum near 18x in affected protocols. Initial insider allocation is uninformative about post-distribution concentration (Pearson $r = 0.05$, $p = 0.76$, $N = 37$), while current insider wallet retention is associated with concentration (Spearman $\rho = 0.48$, $p = 0.003$, $N = 37$); the non-insider HHI check supports this as more than a purely mechanical result.

DePIN protocols are more concentrated than DeFi protocols after correction (mean HHI 0.071 vs 0.031; Mann-Whitney $p = 0.020$; Cohen's $d = 0.94$), but the finding is descriptive and partially attenuates in the smallest multivariate specification. Subsidy dependence does not generalize as a concentration indicator once Livepeer is excluded ($r = 0.07$, $p = 0.77$, $N = 22$). Delegation usually amplifies voting concentration above token holdings in the 13-protocol voting-HHI sample, with ENS, GMX, HNT, and JUP as dispersion exceptions.

The findings span four within-protocol concentration surfaces (post-exclusion holding, insider retention, delegated voting, and vote-escrowed voting) and one cross-protocol surface deferred to future work (institutional delegates participating in multiple protocols' governance simultaneously).

5.6 Contributions

The paper makes four empirical contributions.

First, generalizable PCA exclusion methodology (Section 3.8). The exclusion methodology applies 133 address exclusions controlled by protocols themselves across 38 protocols. Correcting for these addresses changed affected protocols' HHI by a median factor of 2.3x and up to approximately 18x, showing that prior holder-list studies can measure protocol architecture rather than governance control.

Second, allocation-null and insider-retention evidence (Section 4.4). Initial team and investor allocation does not explain post-distribution concentration in the 37-protocol allocation subsample, while insider wallet presence among top holders is associated with concentration. The count-based non-insider HHI check supports the relationship as more than a direct arithmetic artifact; the balance-based version does not survive conventional significance, so the hierarchy of evidence is explicit.

Third, corrected sector contrast (Section 4.3). DePIN governance concentration exceeds DeFi after PCA correction, with robustness to leave-one-out, permutation, and bootstrap checks. The finding is descriptive and should be read with the Model 3 attenuation and small-sample limits in view.

Fourth, delegation-amplification measurement (Section 4.5). Voting-HHI exceeds post-exclusion holding HHI in nine of thirteen protocols with sufficient governance data, while ENS, GMX, HNT, and JUP show that dispersion is possible under different delegate-program or voting-mechanism designs. Curve and Balancer form a separate vote-escrow class with much larger amplification.

5.7 Limitations

The evidence reported here carries five main limitations. First, the analysis is a March 2026 cross-section, so it supports descriptive association rather than causal inference or trajectory claims. Second, the 40-protocol sample is large enough for descriptive comparison but small for fully specified multivariate regression; Model 3 has only 35 observations. Third, PCA classification depends on public labels, bytecode inspection, and documentation quality, with weaker coverage for some Solana-native protocols. Fourth, voting-HHI data are available only for 13 protocols with sufficient governance activity and comparable data surfaces. Fifth, HHI measures concentration of shares, not pivotal voting power under proposal-specific quorum rules.

The attenuated Model 3 sector coefficient ($p = 0.078$) should be read alongside the bootstrap result reported in Section 4.4: the DePIN coefficient bootstrap interval remains above zero $[+0.001, +0.091]$ in the $N = 35$ multivariate sample, but the small sample limits how cleanly sector can be isolated from correlated controls.

The subsidy result is especially sensitive to one observation: excluding Livepeer yields a clean null (full LOO + Spearman + log-transform robustness in Section 4.6), so the main text treats subsidy as a non-generalizing association rather than a stable indicator.

Measurement extensions that would refine these findings without changing the core claim are enumerated in Section 5.8.

5.8 Future Research

Four research directions have highest priority. First, longitudinal governance tracking: monthly snapshots of HHI and Gini across the full sample would convert the current cross-section into a panel dataset, enabling analysis of deconcentration trajectories and the governance effects of specific reform events (e.g., delegation program launches, vesting cliff expirations). A companion panel analysis of quarterly HHI trajectories for 14 governance tokens

([Supplementary File S7](#)) documents that 11 of 14 protocols exhibit monotonic governance HHI decay over 24 months post-token-generation event, with vote-escrow (CRV), reservoir-drip (COMP), and claim-window (ENS) designs as exceptions; this longitudinal pattern is consistent with the cross-sectional allocation null reported in Section 4.4.

Second, sample expansion: the current cross-section of 40 protocols supports descriptive analysis; expanding the sample to 50 to 60 protocols would enable the regression, panel, and event-study analyses currently specified as replication-ready designs.

Third, cross-protocol delegate overlap: the Snapshot-side overlap pattern in Section 4.5.5 should be extended to the Tally delegate registry and Solana governance surfaces to test whether the same professional governance actors concentrate influence across protocols.

Fourth, extending concentration measurement from holding-based HHI to coalition-based power indices (Shapley & Shubik, 1954; Banzhaf, 1968) would quantify whether token holders' pivotal probability diverges from nominal share. Partial-sample computations across five Tally-sourced protocols (Section 4.6; full detail in [Supplementary File S11](#)) show Shapley-Shubik tracks voting-HHI at Pearson $r = 0.999$ under simple-majority quorum, with one protocol-specific Banzhaf divergence (AAVE top-1 Banzhaf 33.5% vs SS 26.7% vs share 21.9%). The 3 Snapshot-sourced protocols (DIMO, LDO, WXM) lack complete delegate-by-delegate weight data and remain the natural expansion target along with proposal-level quorum-configuration data. Recent game-theoretic results (Kiayias et al., 2025) indicate Shapley-based reward distributions achieve bounded Price of Stability ($4/3$) in oceanic staking models while resisting Sybil stake-splitting attacks, supporting the measurement extension's theoretical grounding and computational feasibility for the sample sizes in the current dataset.

B2's forward agenda is empirical: convert the current cross-section into panel evidence, expand the sample, and extend voting-HHI to coalition-based power measures.

Falsifiable forward claims. The cross-sectional design supports four falsifiable predictions about subsequent cross-protocol governance trajectories, each operationalized for future panel-data and event-study work:

1. **Sector persistence:** DePIN protocols launching in 2026-2027 will exhibit higher post-distribution holding HHI than DeFi protocols launching in the same period (Mann-Whitney test on entry-cohort distributions at T+12 and T+24 months).
2. **Delegation amplification universality:** any protocol introducing a formal delegate program will produce voting HHI greater than holding HHI within twelve months of launch, with amplification magnitude varying by delegate-program design rather than by sector.
3. **Insider-retention persistence:** protocols whose insider fraction of top-10 holders increases between baseline and subsequent measurement will exhibit increased holding HHI; protocols whose insider fraction decreases will exhibit decreased holding HHI (within-protocol panel correlation).
4. **Delegate-program design as moderator:** among post-launch protocols in Prediction 2's set, those that adopt broad-community-delegate distribution design choices will exhibit lower amplification ratios than those that do not.

All four predictions inherit the methodology applied in this paper (PCA-symmetric exclusion via the five-class typology; HHI on post-exclusion top-1,000 holders; Mann-Whitney for two-sample tests; bootstrap 95% percentile intervals). Pre-registration commitment: hypothesis specifications, statistical tests, and falsification thresholds will be deposited on the Open Science Framework prior to data collection for any panel-data or event-study extension. Full operationalization detail per prediction (data requirements; explicit falsification thresholds; statistical test specifications) in [Supplementary File S17](#).

5.9 Reproducibility and Data Availability

The 40-protocol governance concentration dataset, exclusions log (133 protocol-controlled address exclusions across 38 protocols), canonical regression dataset including Token Terminal subsidy ratios and on-chain financial data, Hivemapper / io.net / Aethir holder distributions used for Table 3, veCRV cumulative deposit data underlying the Section 4.5.2 ve-locking analysis, and Theil/Atkinson supplementary metrics used in the Section 4.6 robustness check are available in the replication repository linked in the supplementary materials. Replication scripts for each table and figure are provided in Supplementary File S5; the cross-protocol Pearson and Spearman correlations reported across Sections 4.3, 4.4, and 4.6 can be reproduced via the replication scripts using only the linked public-source data. The exclusion methodology is documented at the address-by-address granularity in Supplementary File S6.

6 Conclusion

Governance concentration in token protocols is weakly associated with launch-time allocation and poorly measured by naive holder lists. Across forty protocols, PCA exclusion corrects systematic inflation in prior concentration studies (median factor 2.3x), initial insider allocation is uninformative (Pearson $r = 0.05$, $p = 0.76$, $N = 37$), and post-distribution insider wallet retention is associated with higher concentration (Spearman $\rho = 0.48$, $p = 0.003$, $N = 37$). DePIN protocols exhibit higher concentration than DeFi after measurement correction (Mann-Whitney $p = 0.020$, Cohen's $d = 0.94$), and delegation amplifies voting power above holdings in nine of thirteen protocols with sufficient governance data (ratios 1.6x to 9.9x; mean approximately 4.9x).

For practitioners and researchers, the actionable implication is to audit current holder composition, PCA exclusions, and delegate-program design rather than launch allocation percentages. ENS and Anyone Protocol show that lower concentration and delegation-mediated dispersion are achievable through deliberate institutional design. Longitudinal panel analysis and pre-registered falsifiable predictions are specified in Section 5.8 and Supplementary File S17.

Supplementary Material

The following supplementary files accompany this paper:

S0: Canonical statistics ledger for every B2 main-text statistic.

S1: Metric definitions for holder concentration, voting concentration, subsidy ratios, and related variables.

- S2: Protocol codebook and optional governance-scoring templates retained for researchers extending the broader project. Not load-bearing for B2.
- S3: Optional governance-scoring tables retained as supplementary materials. Not load-bearing for B2.
- S4: Events table for future event-study analyses.
- S5: Empirical pipeline specification, including replication-ready designs for cross-sectional, panel, event-study, and qualitative analyses.
- S6: Source provenance, revenue standardization, raw holder-data workflow, and address-level PCA documentation.
- S7: Quarterly HHI panel for 14 governance tokens across eight quarters post-token-generation event.
- S8: Expanded subsidy-concentration analysis combining Token Terminal revenue data and on-chain emission or fee ratios.
- S9: Theil and Atkinson concentration-metric robustness analysis.
- S10: PCA classification robustness across Specs A through E.
- S11: Shapley-Shubik and Banzhaf power-index calculations.
- S12: PCA-symmetric voting-HHI and signer-side functional PCA robustness.
- S15: Optimism worked example retained as a supplementary case.
- S16: Aethir, IoTeX, and ENS sensitivity analyses.
- S17: Falsifiable predictions and pre-registration specification.

Replication data: https://github.com/Research-Publications-and-Data/Tokenomics-As-Institutional_Design

Acknowledgments

The author thanks Token Terminal for complimentary API access. DeFi protocol subsidy ratios use Token Terminal fee, revenue, and incentive data; DePIN protocol subsidy ratios use on-chain burn and mint data from Dune Analytics due to systematic measurement gaps in existing analytics platforms for hardware network economics.

References

- Birchall, J. (2011). People-centred businesses: Co-operatives, mutuals and the idea of membership. Palgrave Macmillan.
- Hansmann, H. (1996). The ownership of enterprise. Harvard University Press.

- International Cooperative Alliance. (1995). Statement on the cooperative identity. ICA.
- Scholz, T. (2016). Platform cooperativism: Challenging the corporate sharing economy. Rosa Luxemburg Stiftung.
- Alshater, M. M. (2026). Decentralized physical infrastructure networks (DePIN) tokenomics. *Frontiers in Blockchain*, 8, 1644115. <https://doi.org/10.3389/fbloc.2025.1644115>
- Appel, I., & Grennan, J. (2023). Control of decentralized autonomous organizations. *AEA Papers and Proceedings*, 113, 182–185. <https://doi.org/10.1257/pandp.20231119>
- Aramonte, S., Huang, W., & Schrimpf, A. (2021). DeFi risks and the decentralisation illusion. *BIS Quarterly Review*, December 2021, 21–36. https://www.bis.org/publ/qtrpdf/r_qt2112b.htm
- Arnosti, N., & Weinberg, S. M. (2022). Bitcoin: A natural oligopoly. *Management Science*, 68(7), 4755–4771. <https://doi.org/10.1287/mnsc.2021.4095>
- Atkinson, A. B. (1970). On the measurement of inequality. *Journal of Economic Theory*, 2(3), 244–263. [https://doi.org/10.1016/0022-0531\(70\)90039-6](https://doi.org/10.1016/0022-0531(70)90039-6)
- Atzori, M. (2017). Blockchain technology and decentralized governance: Is the state still necessary? *Journal of Governance and Regulation*, 6(1), 45–62. https://doi.org/10.22495/jgr_v6_i1_p5
- Ballandies, M. C., Wang, H., Law, A. C. C., Yang, J. C., Gosken, C., & Andrew, M. (2023). A taxonomy for blockchain-based decentralized physical infrastructure networks (DePIN). 2023 IEEE 9th World Forum on Internet of Things (WF-IoT).
- Bane, D., & Gala, S. (2025). State of DePIN 2024. Messari. <https://messari.io/report/state-of-depin-2024>
- Bane, D., & Gala, S. (2026). State of DePIN 2025. Messari & Escape Velocity. <https://messari.io/report/state-of-depin-2025>
- Barbureau, T., Smethurst, R., Papageorgiou, O., Sedlmeir, J., & Fridgen, G. (2023). Decentralised finance’s timocratic governance: The distribution and exercise of tokenised voting rights. *Technology in Society*, 73, 102251. <https://doi.org/10.1016/j.techsoc.2023.102251>
- Banzhaf, J. F., III. (1968). One man, 3.312 votes: A mathematical analysis of the Electoral College. *Villanova Law Review*, 13(2), 304–332.
- Blockworks. (2026). GEODNET analytics dashboard. Retrieved April 2026, from <https://blockworks.com/analytics/geodnet>
- Blockworks. (2026). Helium financials dashboard. Retrieved April 2026, from <https://blockworks.com/analytics/helium/helium-financials>
- Brennan, G., & Buchanan, J. M. (1985). *The reason of rules: Constitutional political economy*. Cambridge University Press.
- Buchanan, J. M., & Tullock, G. (1962). *The calculus of consent: Logical foundations of constitutional democracy*. University of Michigan Press.

Buterin, V., Conner, E., Dudley, R., Slipper, M., & Norder, I. (2019). EIP-1559: Fee market change for ETH 1.0 chain (Ethereum Improvement Proposal). <https://eips.ethereum.org/EIPS/eip-1559>

Calzada, I. (2025). The political economy of Web3 platformization: Innovation systems, reaching the moon, governing the ghetto. *Digital*, 5(4), 62. <https://doi.org/10.3390/digital5040062>

CoinGecko. (2026). Cryptocurrency market data. Retrieved March 2026, from <https://www.coingecko.com>

Cong, L. W., Li, Y., & Wang, N. (2021). Tokenomics: Dynamic adoption and valuation. *Review of Financial Studies*, 34(3), 1105–1155. <https://doi.org/10.1093/rfs/hhaa089>

Cong, L. W., Rabetti, D., Wang, C. C. Y., & Yan, Y. (2025). Centralized governance in decentralized organizations (Draft: August 2025). Working Paper.

Cohen, J. (1988). *Statistical power analysis for the behavioral sciences* (2nd ed.). Lawrence Erlbaum Associates.

Cohen, J. (1989). Deliberation and democratic legitimacy. In A. Hamlin & P. Pettit (Eds.), *The good polity: Normative analysis of the state* (pp. 17-34). Blackwell.

De Filippi, P., Mannan, M., Cossar, S., Merk, T., & Kamalova, J. (2024). Blockchain technology and polycentric governance (Research Project Report). European University Institute, Robert Schuman Centre. <https://hdl.handle.net/1814/77030>

De Filippi, P., & Wright, A. (2018). *Blockchain and the law: The rule of code*. Harvard University Press. <https://doi.org/10.4159/9780674985933>

DuPont, Q. (2014). The politics of cryptography: Bitcoin and the ordering machines. *Journal of Peer Production*, 4. <https://peerproduction.net>

Dune Analytics. (2026). On-chain governance token holder data [Data set]. Retrieved March 2026, from <https://dune.com>

Eisermann, T., Campajola, C., Tessone, C. J., & Teixeira, A. S. (2025). Concentration in governance control across decentralised finance protocols. *EPJ Data Science*, 14, 85. <https://doi.org/10.1140/epjds/s13688-025-00610-5>

Esposito, M., Tse, T., & Goh, D. (2025). Decentralizing governance: Exploring the dynamics and challenges of digital commons and DAOs. *Frontiers in Blockchain*, 8, 1538227. <https://doi.org/10.3389/fbloc.2025.1538227>

Feichtinger, R., Fritsch, R., Vonlanthen, Y., & Wattenhofer, R. (2023). The hidden shortcomings of (D)AOs: An empirical study of on-chain governance. In *International Conference on Financial Cryptography and Data Security* (pp. 165–185). Springer.

Fritsch, R., Müller, M., & Wattenhofer, R. (2024). Analyzing voting power in decentralized governance: Who controls DAOs? *Blockchain: Research and Applications*, 5(3), 100208. <https://doi.org/10.1016/j.bcra.2024.100208>

- Gersbach, H., Schneider, M., & Shahkar, P. (2024). Effects of vote delegation in blockchains: Who wins? arXiv:2408.05410 (revised July 2025). <https://doi.org/10.48550/arXiv.2408.05410>
- Hall, A. B., & Miyazaki, S. (2024). What happens when anyone can be your representative? Studying the use of liquid democracy for high-stakes decisions in online platforms. Stanford GSB Working Paper. <https://doi.org/10.31219/osf.io/k3hcd>
- Han, J., Lee, J., & Li, T. (2023). DAO governance. European Corporate Governance Institute Finance Working Paper No. 1119/2025. <https://doi.org/10.2139/ssrn.4346581>
- Han, J., Lee, J., & Li, T. (2024). A review of DAO governance: Recent literature and emerging trends. European Corporate Governance Institute Finance Working Paper No. 1044/2025. <https://doi.org/10.2139/ssrn.5074046>
- Hardin, G. (1968). The tragedy of the commons. *Science*, 162(3859), 1243-1248. <https://doi.org/10.1126/science.162.3859.1243>
- Hirschman, A. O. (1964). The paternity of an index. *American Economic Review*, 54(5), 761.
- Holmstrom, B. (1979). Moral hazard and observability. *The Bell Journal of Economics*, 10(1), 74-91. <https://doi.org/10.2307/3003320>
- Howell, S. T., Niessner, M., & Yermack, D. (2020). Initial coin offerings: Financing growth with cryptocurrency token sales. *Review of Financial Studies*, 33(9), 3925–3974. <https://doi.org/10.1093/rfs/hhz131>
- Hurwicz, L. (1973). The design of mechanisms for resource allocation. *American Economic Review*, 63(2), 1–30.
- Hurwicz, L. (1987). Inventing new institutions: The design perspective. *American Journal of Agricultural Economics*, 69(2), 395-402. <https://doi.org/10.2307/1242293>
- Kiayias, A., Koutsoupias, E., Markakis, E., & Tsamopoulos, P. (2025). Pool formation in oceanic games: Shapley value and proportional sharing. In *Proc. 7th Conference on Advances in Financial Technologies (AFT 2025)*, LIPIcs Vol. 354, pp. 21:1–21:24. <https://doi.org/10.4230/LIPIcs.AFT.2025.21>
- Li, S., & Chen, Y. (2024). Governing decentralized autonomous organizations as digital commons. *Journal of Business Venturing Insights*, 22, e00450. <https://doi.org/10.1016/j.jbvi.2024.e00450>
- Michels, R. (1911/1962). *Political parties: A sociological study of the oligarchical tendencies of modern democracy* (E. Paul & C. Paul, Trans.). Free Press. (Original work published 1911)
- Myerson, R. B. (1981). Optimal auction design. *Mathematics of Operations Research*, 6(1), 58–73. <https://doi.org/10.1287/moor.6.1.58>
- Nabben, K., Wang, H., & Zargham, M. (2024). Decentralised governance for autonomous cyber-physical systems. arXiv preprint, arXiv:2408.16764. <https://arxiv.org/abs/2408.16764>

- Nadler, M., & Schär, F. (2020). Decentralized finance, centralized ownership? An iterative mapping process to measure protocol token distribution. arXiv:2012.09306. <https://doi.org/10.48550/arXiv.2012.09306>
- North, D. C. (1990). Institutions, institutional change and economic performance. Cambridge University Press.
- Olson, M. (1965). The logic of collective action: Public goods and the theory of groups. Harvard University Press.
- Ostrom, E. (1990). Governing the commons: The evolution of institutions for collective action. Cambridge University Press.
- Optimism. (2024). Optimism governance documentation: Citizens' House. <https://community.optimism.io/docs/governance/>
- Polanyi, K. (1944). The great transformation: The political and economic origins of our time. Farrar & Rinehart.
- Rikken, O., Janssen, M., & Kwee, Z. (2023). Governance impacts of blockchain-based decentralized autonomous organizations: An empirical analysis. Policy Design and Practice, 6(4), 465–487. <https://doi.org/10.1080/25741292.2023.2270220>
- Roughgarden, T. (2020). Transaction fee mechanism design for the Ethereum blockchain: An economic analysis of EIP-1559. arXiv:2012.00854. <https://arxiv.org/abs/2012.00854>
- Roughgarden, T. (2021). Transaction fee mechanism design. ACM SIGecom Exchanges, 19(1), 52–55. <https://doi.org/10.1145/3476436.3476445>
- Sai, A. R., Buckley, J., & Le Gear, A. (2021). Characterizing wealth inequality in cryptocurrencies. Frontiers in Blockchain, 4, 730122. <https://doi.org/10.3389/fbloc.2021.730122>
- Sawinyh, N. (2025, October 29). Points-based distribution programs in Web3: A comprehensive guide. DeFiPrime. <https://defiprime.com/points-based-token-distribution-programs-web3>
- Schär, F. (2021). Decentralized finance: On blockchain- and smart contract-based financial markets. Federal Reserve Bank of St. Louis Review, 103(2), 153–174. <https://doi.org/10.20955/r.103.153-74>
- Schneider, N. (2024). Governable spaces: Democratic design for online life. University of California Press (Luminos Open Access). <https://doi.org/10.1525/luminos.181>
- Shapley, L. S. (1953). A value for n-person games. In H. W. Kuhn & A. W. Tucker (Eds.), Contributions to the theory of games, Volume II (Annals of Mathematics Studies 28, pp. 307–317). Princeton University Press.
- Shapley, L. S., & Shubik, M. (1954). A method for evaluating the distribution of power in a committee system. American Political Science Review, 48(3), 787-792. <https://doi.org/10.2307/1951053>
- Theil, H. (1967). Economics and information theory. North-Holland.

Tirole, J. (2001). Corporate governance. *Econometrica*, 69(1), 1-35.
<https://doi.org/10.1111/1468-0262.00177>

Tchuente, G. (2025). Too big to monitor? Network scale and the breakdown of decentralized monitoring. *arXiv:2511.23320*. <https://doi.org/10.48550/arXiv.2511.23320>

Van Vulpen, L., & Jansen, S. (2023). Decentralized autonomous organization design for the commons and the common good. *Frontiers in Blockchain*, 6, 1257313.
<https://doi.org/10.3389/fbloc.2023.1257313>

Voshmgir, S. (2020). *Token economy: How the web3 reinvents the internet* (2nd ed.). Token Kitchen.

Weidener, L., Anwar, A., Richter, T., Pournader, P., & Kartal, A. B. (2025). Delegated voting in decentralized autonomous organizations: A scoping review. *Frontiers in Blockchain*, 8, 1533405.
<https://doi.org/10.3389/fbloc.2025.1533405>

Werner, S. M., Perez, D., Gudgeon, L., Klages-Mundt, A., Harz, D., & Knottenbelt, W. J. (2022). SoK: Decentralized finance (DeFi). In *Proceedings of the 4th ACM Conference on Advances in Financial Technologies* (pp. 30–46). ACM.
<https://doi.org/10.1145/3558535.3559780>

Williamson, O. E. (1999). Public and private bureaucracies: A transaction cost economics perspective. *Journal of Law, Economics, and Organization*, 15(1), 306–342.

Zukowski, Z. (2026a). Who burns the tokens? Fiscal sustainability and demand concentration in GeoDePIN networks. SSRN Working Paper. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=6483619

Zukowski, Z. (2026b). Adaptive Tokenomics: A Systems Engineering Approach to Programmable Incentive Design. SSRN Working Paper No. 6364158.
https://papers.ssrn.com/sol3/papers.cfm?abstract_id=6364158

Zukowski, Z. (2026c). Routing the Dollar: Gateway infrastructure, monetary policy transmission, and the Dollar's international functions in digital markets. SSRN Working Paper No. 6267698. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=6267698

Zukowski, Z. (2026d). Seven Dollars: Legal Structure, Operator Risk, and Diligence Divergence Across Dollar-Denominated Digital Products. SSRN Working Paper No. 6352118.
https://papers.ssrn.com/sol3/papers.cfm?abstract_id=6352118

Zukowski, Z. (2026e). Why token governance concentrates: A Polanyian extension to the fourth fictitious commodity. Working manuscript.